

# FPGA Development for Radar, Radio-Astronomy and Communications

THE  
**RADAR**  
MASTERS COURSE



Dept. Electrical Engineering, University of Cape Town  
Private Bag, Rondebosch, 7701, South Africa  
<http://www.rrsg.uct.ac.za>



Presented by John-Philip Taylor

Convened by Dr Stephen Paine

Day 4 – 2 May 2022

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration

Caching Systems

Flow Control

Practical – NCO and DDC

Numerically-controlled Oscillator



# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration

Caching Systems

Flow Control

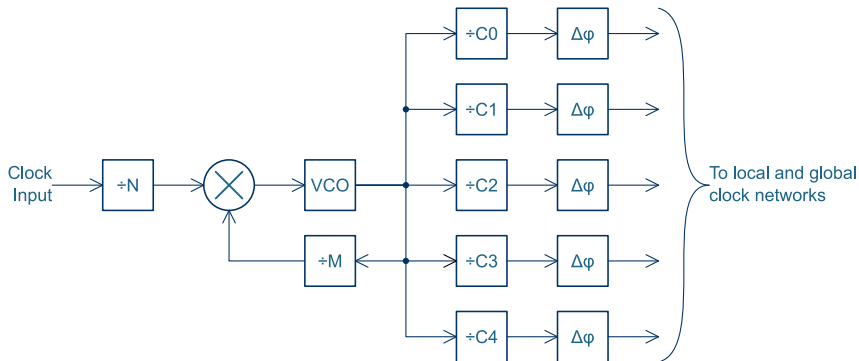
Practical – NCO and DDC


Numerically-controlled Oscillator



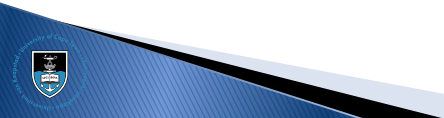
# Phase-locked Loop


2 of 34






- ▶ Useful for generating a range of related clocks
- ▶ Output clock rising and falling edges can be set, independently, to a resolution equal to the VCO period, which is typically about  $1 \text{ ns} \pm 500 \text{ ps}$
- ▶ All output clocks can be considered to be in the same clock domain
- ▶ Noise sensitive applications should avoid PLLs due to high phase noise






- ▶ Useful for generating a range of related clocks
- ▶ Output clock rising and falling edges can be set, independently, to a resolution equal to the VCO period, which is typically about  $1 \text{ ns} \pm 500 \text{ ps}$
- ▶ All output clocks can be considered to be in the same clock domain
- ▶ Noise sensitive applications should avoid PLLs due to high phase noise




- ▶ Useful for generating a range of related clocks
- ▶ Output clock rising and falling edges can be set, independently, to a resolution equal to the VCO period, which is typically about  $1 \text{ ns} \pm 500 \text{ ps}$
- ▶ All output clocks can be considered to be in the same clock domain
- ▶ Noise sensitive applications should avoid PLLs due to high phase noise




- ▶ Useful for generating a range of related clocks
- ▶ Output clock rising and falling edges can be set, independently, to a resolution equal to the VCO period, which is typically about  $1 \text{ ns} \pm 500 \text{ ps}$
- ▶ All output clocks can be considered to be in the same clock domain
- ▶ Noise sensitive applications should avoid PLLs due to high phase noise







- ▶ **Direct:** No compensation, which results in better jitter performance
- ▶ **Normal:** The clock at the register is phase-aligned with the input clock pin
- ▶ **Source Synchronous:** The PLL ensures that the data delay from pin to register and the apparent clock delay from pin to register is the same
- ▶ **Zero-Delay Buffer:** The output clock pin is phase-aligned with the input clock pin



- ▶ **Direct:** No compensation, which results in better jitter performance
- ▶ **Normal:** The clock at the register is phase-aligned with the input clock pin
- ▶ **Source Synchronous:** The PLL ensures that the data delay from pin to register and the apparent clock delay from pin to register is the same
- ▶ **Zero-Delay Buffer:** The output clock pin is phase-aligned with the input clock pin



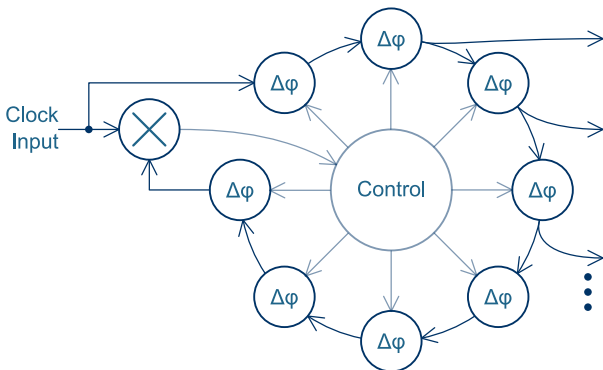
- ▶ **Direct:** No compensation, which results in better jitter performance
- ▶ **Normal:** The clock at the register is phase-aligned with the input clock pin
- ▶ **Source Synchronous:** The PLL ensures that the data delay from pin to register and the apparent clock delay from pin to register is the same
- ▶ **Zero-Delay Buffer:** The output clock pin is phase-aligned with the input clock pin



- ▶ **Direct:** No compensation, which results in better jitter performance
- ▶ **Normal:** The clock at the register is phase-aligned with the input clock pin
- ▶ **Source Synchronous:** The PLL ensures that the data delay from pin to register and the apparent clock delay from pin to register is the same
- ▶ **Zero-Delay Buffer:** The output clock pin is phase-aligned with the input clock pin


# Delay-locked Loop

4 of 34

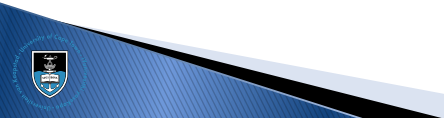


# Delay-locked Loop

4 of 34




- ▶ Useful for performing phase compensation when interfacing to fast peripherals
- ▶ All output clocks can be considered to be in the same clock domain as the input clock



# Delay-locked Loop


4 of 34




- ▶ Useful for performing phase compensation when interfacing to fast peripherals
- ▶ All output clocks can be considered to be in the same clock domain as the input clock

# Clock-Enable Generated


5 of 34









- ▶ Useful when mixing fast and slow systems
- ▶ Potential to reduce overall device power usage
- ▶ Cannot be used for clocking external devices
- ▶ The multi-cycle timing requirements must be manually specified



- ▶ Useful when mixing fast and slow systems
- ▶ Potential to reduce overall device power usage
- ▶ Cannot be used for clocking external devices
- ▶ The multi-cycle timing requirements must be manually specified



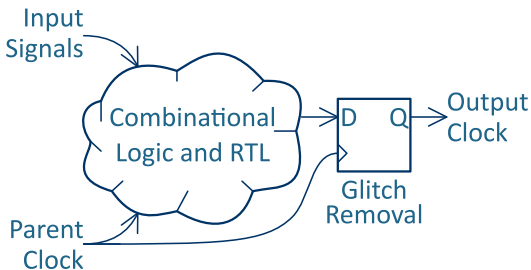
- ▶ Useful when mixing fast and slow systems
- ▶ Potential to reduce overall device power usage
- ▶ Cannot be used for clocking external devices
- ▶ The multi-cycle timing requirements must be manually specified



- ▶ Useful when mixing fast and slow systems
- ▶ Potential to reduce overall device power usage
- ▶ Cannot be used for clocking external devices
- ▶ The multi-cycle timing requirements must be manually specified


# Ripple and Gated Clocks

6 of 34




# Ripple and Gated Clocks


6 of 34



- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.




- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.




- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.






- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.






- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.



- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.






- ▶ Unrelated to all other clocks in the system, including the parent clock
- ▶ Often necessitates complex clock-domain crossing schemes to the general system clock
- ▶ Useful when:
  - ▶ Requiring very low power usage
  - ▶ The frequency must change often
  - ▶ The clock must drive an external device
  - ▶ etc.

# Clock Regions

7 of 34





- ▶ FPGAs are generally organised into regions, each with its own PLL and I/O bank
- ▶ The local clock network is more efficient (and faster) than the global network
- ▶ When laying out the PCB, keep fast I/O in the same region as their clock



- ▶ FPGAs are generally organised into regions, each with its own PLL and I/O bank
- ▶ The local clock network is more efficient (and faster) than the global network
- ▶ When laying out the PCB, keep fast I/O in the same region as their clock



- ▶ FPGAs are generally organised into regions, each with its own PLL and I/O bank
- ▶ The local clock network is more efficient (and faster) than the global network
- ▶ When laying out the PCB, keep fast I/O in the same region as their clock



# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration

Caching Systems

Flow Control


Practical – NCO and DDC

Numerically-controlled Oscillator




```
module RS232(  
    input  ipClk, ipReset,  
    input  ipRS232_Rx, // Asynchronous external signal  
    output [7:0]opData  
);  
  
... // Definitions, state machine boilerplate, etc.  
  
Idle: begin  
    if(!ipRS232_Rx) begin  
        State <= Receiving;  
    end  
end  
  
Receiving: begin  
...  
end
```






# Common Mistake



8 of 34






# Common Mistake


9 of 34



# Register Chain


10 of 34






- ▶ Essential for crossing asynchronous external signals
- ▶ Only useful for crossing one bit at a time (unless the system is tolerant to temporary errors)
- ▶ Often used to cross hand-shaking signals
- ▶ Works in both directions (fast to slow / slow to fast)






- ▶ Essential for crossing asynchronous external signals
- ▶ Only useful for crossing one bit at a time (unless the system is tolerant to temporary errors)
- ▶ Often used to cross hand-shaking signals
- ▶ Works in both directions (fast to slow / slow to fast)



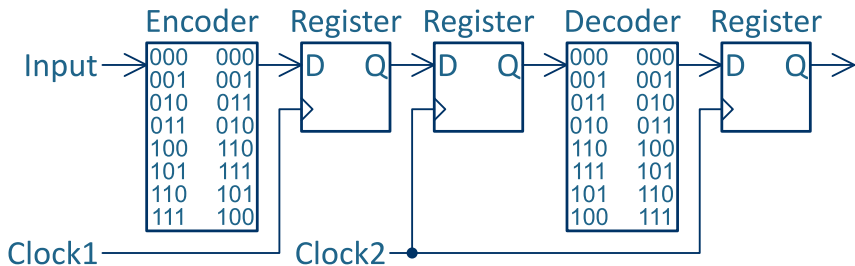
- ▶ Essential for crossing asynchronous external signals
- ▶ Only useful for crossing one bit at a time (unless the system is tolerant to temporary errors)
- ▶ Often used to cross hand-shaking signals
- ▶ Works in both directions (fast to slow / slow to fast)



- ▶ Essential for crossing asynchronous external signals
- ▶ Only useful for crossing one bit at a time (unless the system is tolerant to temporary errors)
- ▶ Often used to cross hand-shaking signals
- ▶ Works in both directions (fast to slow / slow to fast)


# Gray Coding

11 of 34



# Gray Coding


11 of 34



- ▶ Encoded such that only one bit changes at a time
- ▶ Only works for counter data  
(FIFO queue pointers, rotary encoders, etc.)
- ▶ Works in both directions (fast to slow / slow to fast)

# Gray Coding


11 of 34



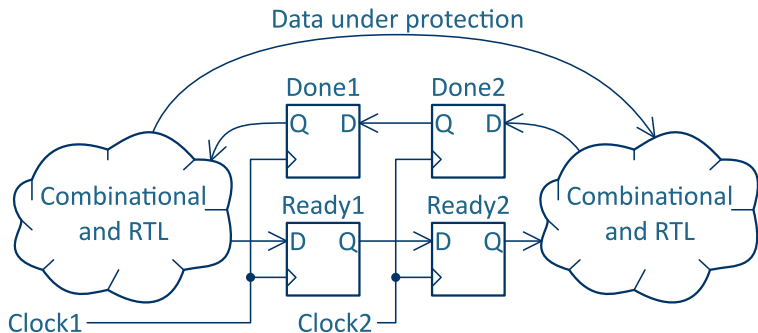
- ▶ Encoded such that only one bit changes at a time
- ▶ Only works for counter data  
(FIFO queue pointers, rotary encoders, etc.)
- ▶ Works in both directions (fast to slow / slow to fast)

# Gray Coding

11 of 34





- ▶ Encoded such that only one bit changes at a time
- ▶ Only works for counter data  
(FIFO queue pointers, rotary encoders, etc.)
- ▶ Works in both directions (fast to slow / slow to fast)







# Hand-Shaking

12 of 34







- ▶ Source must guarantee data stability while the handshaking is taking place
- ▶ Not time-efficient: need to wait for the entire transaction sequence to finish before starting the next one
- ▶ Use only for data that does not need to be crossed often
- ▶ Works in both directions (fast to slow / slow to fast)



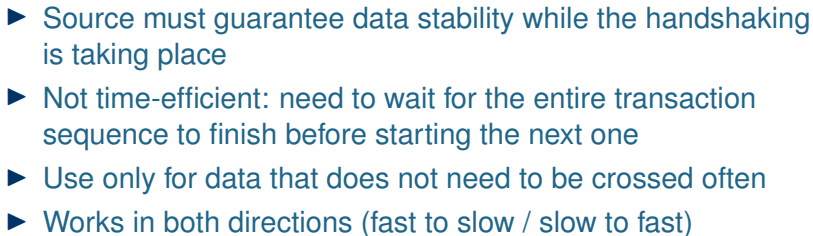
- ▶ Source must guarantee data stability while the handshaking is taking place
- ▶ Not time-efficient: need to wait for the entire transaction sequence to finish before starting the next one
- ▶ Use only for data that does not need to be crossed often
- ▶ Works in both directions (fast to slow / slow to fast)




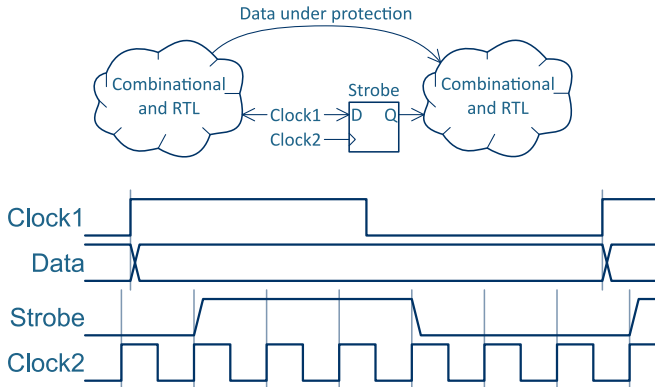
- ▶ Source must guarantee data stability while the handshaking is taking place
- ▶ Not time-efficient: need to wait for the entire transaction sequence to finish before starting the next one
- ▶ Use only for data that does not need to be crossed often
- ▶ Works in both directions (fast to slow / slow to fast)




## 12 of 34










- ▶ The data is valid for as long as the strobe is high
- ▶ As fast as the source clock
- ▶ The strobe edges can be used to perform flow-control
- ▶ Only works when the destination clock is much faster than the source clock
- ▶ Note that the source “clock” can be a strobe signal generated by the source state machine...









- ▶ The data is valid for as long as the strobe is high
- ▶ As fast as the source clock
- ▶ The strobe edges can be used to perform flow-control
- ▶ Only works when the destination clock is much faster than the source clock
- ▶ Note that the source “clock” can be a strobe signal generated by the source state machine...






- ▶ The data is valid for as long as the strobe is high
- ▶ As fast as the source clock
- ▶ The strobe edges can be used to perform flow-control
- ▶ Only works when the destination clock is much faster than the source clock
- ▶ Note that the source “clock” can be a strobe signal generated by the source state machine...





- ▶ The data is valid for as long as the strobe is high
- ▶ As fast as the source clock
- ▶ The strobe edges can be used to perform flow-control
- ▶ Only works when the destination clock is much faster than the source clock
- ▶ Note that the source “clock” can be a strobe signal generated by the source state machine...





- ▶ The data is valid for as long as the strobe is high
- ▶ As fast as the source clock
- ▶ The strobe edges can be used to perform flow-control
- ▶ Only works when the destination clock is much faster than the source clock
- ▶ Note that the source “clock” can be a strobe signal generated by the source state machine...





- ▶ Flow-control signals are crossed separately
- ▶ Control signal crossing latency can be hidden in the length of the queue
- ▶ Mixed-width RAM can be used to keep the data-rate constant across different frequencies (especially useful in DDR-based external memory)





- ▶ Flow-control signals are crossed separately
- ▶ Control signal crossing latency can be hidden in the length of the queue
- ▶ Mixed-width RAM can be used to keep the data-rate constant across different frequencies (especially useful in DDR-based external memory)





- ▶ Flow-control signals are crossed separately
- ▶ Control signal crossing latency can be hidden in the length of the queue
- ▶ Mixed-width RAM can be used to keep the data-rate constant across different frequencies (especially useful in DDR-based external memory)





- ▶ Flow-control signals are crossed separately
- ▶ Control signal crossing latency can be hidden in the length of the queue
- ▶ Mixed-width RAM can be used to keep the data-rate constant across different frequencies (especially useful in DDR-based external memory)





- ▶ Unrelated clock domains must be kept separate in the timing constraints as well
- ▶ Remember clock groups from Day 2?



- ▶ Unrelated clock domains must be kept separate in the timing constraints as well
- ▶ Remember clock groups from Day 2?

```
set_clock_groups -asynchronous \  
-group [get_clocks ADC_Clk]      \  
-group [get_clocks Clk1]         \  
-group [get_clocks Clk2]         \  
-group [get_clocks {SRAM_CLK *altpll_0*}]
```



# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration


Caching Systems

Flow Control


Practical – NCO and DDC

Numerically-controlled Oscillator







- ▶ Often many modules need access to the same resource, but only one module can interface with it at a time
- ▶ Examples include:
  - ▶ External memory (SDRAM, SD-card, etc.)
  - ▶ I<sup>2</sup>C bus
  - ▶ etc.




- ▶ Often many modules need access to the same resource, but only one module can interface with it at a time
- ▶ Examples include:
  - ▶ External memory (SDRAM, SD-card, etc.)
  - ▶ I<sup>2</sup>C bus
  - ▶ etc.



- ▶ Uses “Request” and “Grant” control lines
- ▶ The module would raise a request and then wait until it has been granted access before using the resource
- ▶ All the request lines go to a central control module that administers the granting of the resource




- ▶ Uses “Request” and “Grant” control lines
- ▶ The module would raise a request and then wait until it has been granted access before using the resource
- ▶ All the request lines go to a central control module that administers the granting of the resource



- ▶ Uses “Request” and “Grant” control lines
- ▶ The module would raise a request and then wait until it has been granted access before using the resource
- ▶ All the request lines go to a central control module that administers the granting of the resource






- ▶ Uses “Request” and “Grant” control lines
- ▶ The module would raise a request and then wait until it has been granted access before using the resource
- ▶ All the request lines go to a central control module that administers the granting of the resource

# Round-robin Mutual Exclusion

18 of 34




- ▶ The requests are serviced in a circular order
- ▶ Guaranteed no starvation
- ▶ Does not scale well (in the case of many modules, many clock-cycles must be wasted checking each module's request line)



# Round-robin Mutual Exclusion

18 of 34




- ▶ The requests are serviced in a circular order
- ▶ Guaranteed no starvation
- ▶ Does not scale well (in the case of many modules, many clock-cycles must be wasted checking each module's request line)



# Round-robin Mutual Exclusion

18 of 34




- ▶ The requests are serviced in a circular order
- ▶ Guaranteed no starvation
- ▶ Does not scale well (in the case of many modules, many clock-cycles must be wasted checking each module's request line)



# Round-robin Mutual Exclusion


18 of 34



- ▶ The requests are serviced in a circular order
- ▶ Guaranteed no starvation
- ▶ Does not scale well (in the case of many modules, many clock-cycles must be wasted checking each module's request line)

# Priority-based Mutual Exclusion

19 of 34




- ▶ When more than one module requests access, the one with higher priority always receives the grant
- ▶ Can result in starvation of low-priority modules
- ▶ Fast (implementation can be a combinational circuit, providing a response within the same clock cycle)
- ▶ Scales well



# Priority-based Mutual Exclusion

19 of 34




- ▶ When more than one module requests access, the one with higher priority always receives the grant
- ▶ Can result in starvation of low-priority modules
- ▶ Fast (implementation can be a combinational circuit, providing a response within the same clock cycle)
- ▶ Scales well



# Priority-based Mutual Exclusion

19 of 34




- ▶ When more than one module requests access, the one with higher priority always receives the grant
- ▶ Can result in starvation of low-priority modules
- ▶ Fast (implementation can be a combinational circuit, providing a response within the same clock cycle)
- ▶ Scales well





# Priority-based Mutual Exclusion

19 of 34




- ▶ When more than one module requests access, the one with higher priority always receives the grant
- ▶ Can result in starvation of low-priority modules
- ▶ Fast (implementation can be a combinational circuit, providing a response within the same clock cycle)
- ▶ Scales well





# Priority-based Mutual Exclusion

19 of 34




- ▶ When more than one module requests access, the one with higher priority always receives the grant
- ▶ Can result in starvation of low-priority modules
- ▶ Fast (implementation can be a combinational circuit, providing a response within the same clock cycle)
- ▶ Scales well






- ▶ The arbiter makes it look as if the resource has multiple independent interfaces. There are no control lines other than the interface itself.
- ▶ Often implemented by means of an embedded mutual exclusion unit (round-robin or priority based)
- ▶ The I<sup>2</sup>C standard implements arbitration by collision-detection and random back-off: the modules monitor their own output, and when there is a mismatch, the module waits and tries again later.




- ▶ The arbiter makes it look as if the resource has multiple independent interfaces. There are no control lines other than the interface itself.
- ▶ Often implemented by means of an embedded mutual exclusion unit (round-robin or priority based)
- ▶ The I<sup>2</sup>C standard implements arbitration by collision-detection and random back-off: the modules monitor their own output, and when there is a mismatch, the module waits and tries again later.



- ▶ The arbiter makes it look as if the resource has multiple independent interfaces. There are no control lines other than the interface itself.
- ▶ Often implemented by means of an embedded mutual exclusion unit (round-robin or priority based)
- ▶ The I<sup>2</sup>C standard implements arbitration by collision-detection and random back-off: the modules monitor their own output, and when there is a mismatch, the module waits and tries again later.

# Avalon Bus Arbitration

21 of 34



# Coffee Break...

---

22 of 34





# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration


Caching Systems

Flow Control


Practical – NCO and DDC

Numerically-controlled Oscillator







- ▶ Lines are typically one page in size (4 kiB in modern PCs, typically smaller on embedded systems)
- ▶ The controller does address translation between the external and cache line addresses
- ▶ Read is on-demand, flushing the least recently used line and loading the content into it
- ▶ Controller blocks the user module when a cache miss occurs



- ▶ Lines are typically one page in size (4 kiB in modern PCs, typically smaller on embedded systems)
- ▶ The controller does address translation between the external and cache line addresses
- ▶ Read is on-demand, flushing the least recently used line and loading the content into it
- ▶ Controller blocks the user module when a cache miss occurs



- ▶ Lines are typically one page in size (4 kiB in modern PCs, typically smaller on embedded systems)
- ▶ The controller does address translation between the external and cache line addresses
- ▶ Read is on-demand, flushing the least recently used line and loading the content into it
- ▶ Controller blocks the user module when a cache miss occurs



- ▶ Lines are typically one page in size (4 kiB in modern PCs, typically smaller on embedded systems)
- ▶ The controller does address translation between the external and cache line addresses
- ▶ Read is on-demand, flushing the least recently used line and loading the content into it
- ▶ Controller blocks the user module when a cache miss occurs



- ▶ Reading is continuous, and the user module is only blocked when the queue is empty
- ▶ The queue is typically multiple pages in size
- ▶ New data is fetched as soon as there is one page open on the queue




- ▶ Reading is continuous, and the user module is only blocked when the queue is empty
- ▶ The queue is typically multiple pages in size
- ▶ New data is fetched as soon as there is one page open on the queue




- ▶ Reading is continuous, and the user module is only blocked when the queue is empty
- ▶ The queue is typically multiple pages in size
- ▶ New data is fetched as soon as there is one page open on the queue







- ▶ Writing is continuous, and the user module is only blocked when the queue is full
- ▶ Alternatively, the user module is never blocked, and data is lost when the queue is full
- ▶ The flush unit automatically flushes the queue when there is one page of data in it
- ▶ The user module can also manually flush the queue if required




- ▶ Writing is continuous, and the user module is only blocked when the queue is full
- ▶ Alternatively, the user module is never blocked, and data is lost when the queue is full
- ▶ The flush unit automatically flushes the queue when there is one page of data in it
- ▶ The user module can also manually flush the queue if required




- ▶ Writing is continuous, and the user module is only blocked when the queue is full
- ▶ Alternatively, the user module is never blocked, and data is lost when the queue is full
- ▶ The flush unit automatically flushes the queue when there is one page of data in it
- ▶ The user module can also manually flush the queue if required




- ▶ Writing is continuous, and the user module is only blocked when the queue is full
- ▶ Alternatively, the user module is never blocked, and data is lost when the queue is full
- ▶ The flush unit automatically flushes the queue when there is one page of data in it
- ▶ The user module can also manually flush the queue if required




- ▶ Two independent cache system can potentially refer to the same place in external storage
- ▶ When the one user module reads what the other is writing, the two cache system must remain up to date with each other
- ▶ When the two user modules write to different places in the same page, a flush must not overwrite the other's data
- ▶ Many schemes exist to mitigate these and other issues, which are outside the scope of this course



- ▶ Two independent cache system can potentially refer to the same place in external storage
- ▶ When the one user module reads what the other is writing, the two cache system must remain up to date with each other
- ▶ When the two user modules write to different places in the same page, a flush must not overwrite the other's data
- ▶ Many schemes exist to mitigate these and other issues, which are outside the scope of this course



- ▶ Two independent cache system can potentially refer to the same place in external storage
- ▶ When the one user module reads what the other is writing, the two cache system must remain up to date with each other
- ▶ When the two user modules write to different places in the same page, a flush must not overwrite the other's data
- ▶ Many schemes exist to mitigate these and other issues, which are outside the scope of this course



- ▶ Two independent cache system can potentially refer to the same place in external storage
- ▶ When the one user module reads what the other is writing, the two cache system must remain up to date with each other
- ▶ When the two user modules write to different places in the same page, a flush must not overwrite the other's data
- ▶ Many schemes exist to mitigate these and other issues, which are outside the scope of this course



# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration

Caching Systems

Flow Control

Practical – NCO and DDC

Numerically-controlled Oscillator





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss





- ▶ No feedback from the destination to the source
- ▶ Assumes that the destination will handle the data, or
- ▶ The application can accept data loss







- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination





- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination





- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination






- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination






- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination





- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination





- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination





- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination







- ▶ The source sends one packet at a time
- ▶ And waits for the response from the destination





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols







- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The destination advertises a window
- ▶ The source aims to fill the window with data
- ▶ This is used in TCP, among other protocols





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times







- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times







- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ The source assumes that the destination has a certain space
- ▶ The destination acknowledges all incoming packets
- ▶ The source sends packets aiming to keep the credit on zero
- ▶ Works best when the credit is large enough to keep both directions full at all times





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols



## Credit Based (2)

31 of 34



- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols







- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols







- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols





- ▶ Instead of the source starting with credit, an alternative algorithm lets the destination provide credit to the source
- ▶ This is used in PCIe, among other protocols



# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration

Caching Systems

Flow Control

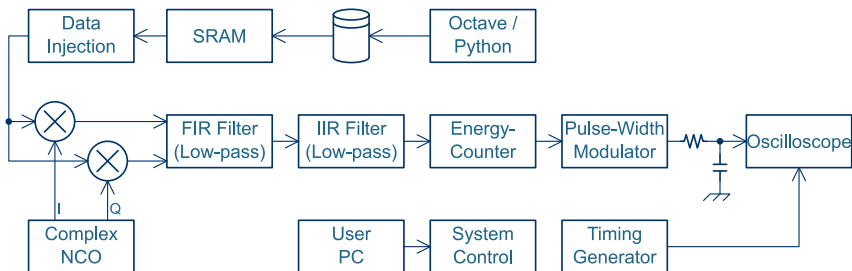
Practical – NCO and DDC

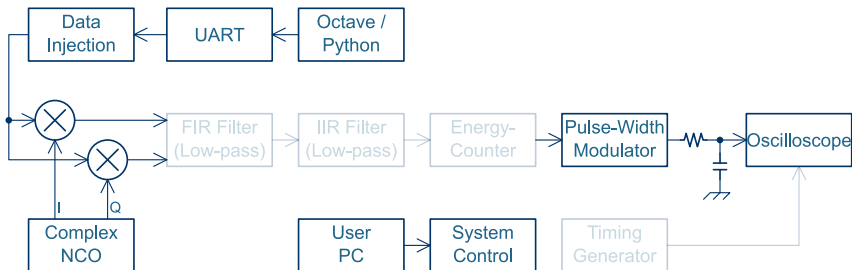
Numerically-controlled Oscillator



# Practical – NCO and DDC

32 of 34





# Outline

---

Advanced Clocking

Clock Domains

Mutual Exclusion and Arbitration


Caching Systems

Flow Control

Practical – NCO and DDC


Numerically-controlled Oscillator






- ▶ Integrate frequency to obtain phase
- ▶ Use on-chip memory blocks as a look-up table to convert phase to the intended waveform
- ▶ Use dual-port ROM to obtain sine and cosine from the same LUT
- ▶ Higher resolution can be obtained with a CORDIC-based sine-cosine module







- Integrate frequency to obtain phase
- Use on-chip memory blocks as a look-up table to convert phase to the intended waveform
- Use dual-port ROM to obtain sine and cosine from the same LUT
- Higher resolution can be obtained with a CORDIC-based sine-cosine module




- ▶ Integrate frequency to obtain phase
- ▶ Use on-chip memory blocks as a look-up table to convert phase to the intended waveform
- ▶ Use dual-port ROM to obtain sine and cosine from the same LUT
- ▶ Higher resolution can be obtained with a CORDIC-based sine-cosine module



- ▶ Integrate frequency to obtain phase
- ▶ Use on-chip memory blocks as a look-up table to convert phase to the intended waveform
- ▶ Use dual-port ROM to obtain sine and cosine from the same LUT
- ▶ Higher resolution can be obtained with a CORDIC-based sine-cosine module



- ▶ Integrate frequency to obtain phase
- ▶ Use on-chip memory blocks as a look-up table to convert phase to the intended waveform
- ▶ Use dual-port ROM to obtain sine and cosine from the same LUT
- ▶ Higher resolution can be obtained with a CORDIC-based sine-cosine module



$$f_{out} = f_s \cdot \frac{f_{in}}{2^N}, \quad N = 32, \quad f_s = 50 \text{ MHz}$$



# Select References

34 of 34



Stephen Brown and Zvonko Vranesic  
Fundamentals of Digital Logic with Verilog Design, 2<sup>nd</sup> Edition  
ISBN 978-0-07-721164-6



Merrill L Skolnik  
Introduction to RADAR Systems  
ISBN 978-0-07-288138-7



Mark A. Richards and James A. Scheer  
Principles of Modern Radar: Basic Principles  
ISBN 978-1-89-112152-4



Deepak Kumar Tala  
World of ASIC  
<http://www.asic-world.com/>



Jean P. Nicolle  
FPGA 4 Fun  
<http://www.fpga4fun.com/>



# FPGA Development for Radar, Radio-Astronomy and Communications

THE  
**RADAR**  
MASTERS COURSE



Dept. Electrical Engineering, University of Cape Town  
Private Bag, Rondebosch, 7701, South Africa  
<http://www.rrsg.uct.ac.za>



Presented by John-Philip Taylor

Convened by Dr Stephen Paine

Day 4 – 2 May 2022