

# FPGA Development for Radar, Radio-Astronomy and Communications

THE  
**RADAR**  
MASTERS COURSE



Dept. Electrical Engineering, University of Cape Town  
Private Bag, Rondebosch, 7701, South Africa  
<http://www.rrsg.uct.ac.za>



Presented by John-Philip Taylor  
Convened by Dr Stephen Paine

Day 6 – 4 May 2022

Advanced Timing Constraints

FIR Filters

Practical – FIR Filter

Projects

Tips and Tricks

Conclusion



# Outline

---

Advanced Timing Constraints

FIR Filters

Practical – FIR Filter

Projects

Tips and Tricks

Conclusion



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)





- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)

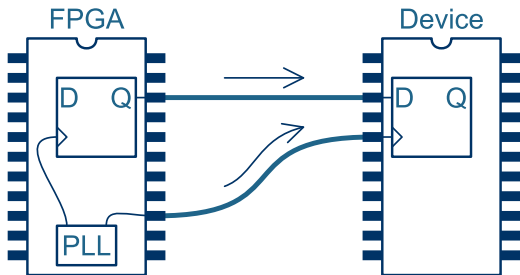


- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)



- ▶ Synopsis Design Constraints start with the ideal case:
  - ▶ There are no PCB trace delays
  - ▶ The external setup requirement is zero
  - ▶ The external hold requirement is zero
  - ▶ The external propagation delay is zero
- ▶ Increase the maximum output delay for external setup timing
- ▶ Decrease the minimum output delay for external hold timing
- ▶ Specify the minimum and maximum input delays according to the external propagation delay parameters
- ▶ Worsen the situation with PCB trace delays and uncertainties (clock jitter, manufacturing tolerances, etc.)

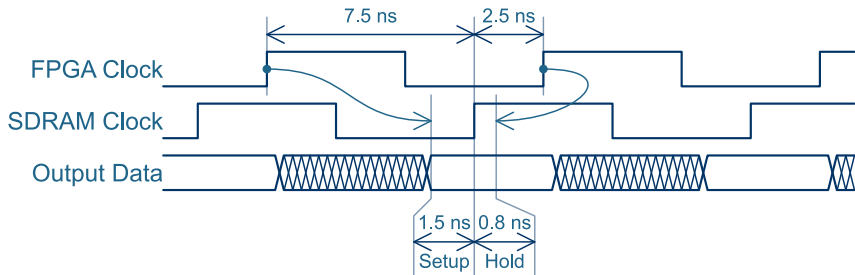




- ▶ FPGA internal delays and PCB trace delays cancel
- ▶ The important parameters are:
  - ▶ Setup and hold times of the external device
  - ▶ Clock jitter and other uncertainties
- ▶ Shift the external clock to ease the hold margin

# External Timing – Output

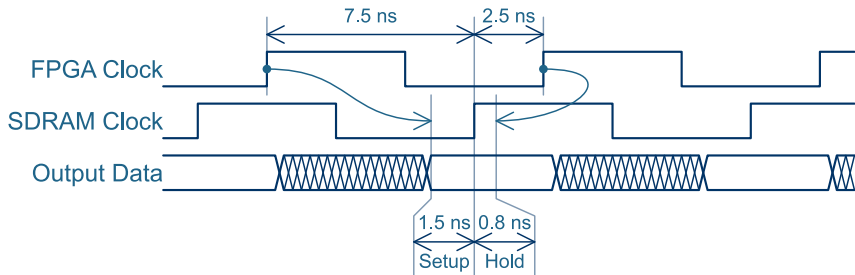
3 of 40



- ▶ FPGA internal delays and PCB trace delays cancel
- ▶ The important parameters are:
  - ▶ Setup and hold times of the external device
  - ▶ Clock jitter and other uncertainties
- ▶ Shift the external clock to ease the hold margin

# External Timing – Output

3 of 40



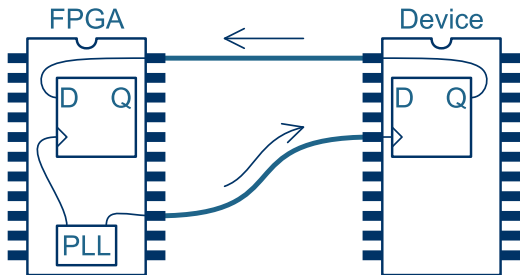
# Suppose 100 ps uncertainty

```
set_output_delay -max -clock DRAM_CLK 1.6 [get_ports opDRAM*]  
set_output_delay -min -clock DRAM_CLK -0.9 [get_ports opDRAM*]  
set_output_delay -max -clock DRAM_CLK 1.6 [get_ports bpDRAM*]  
set_output_delay -min -clock DRAM_CLK -0.9 [get_ports bpDRAM*]
```



# External Timing – Input

4 of 40

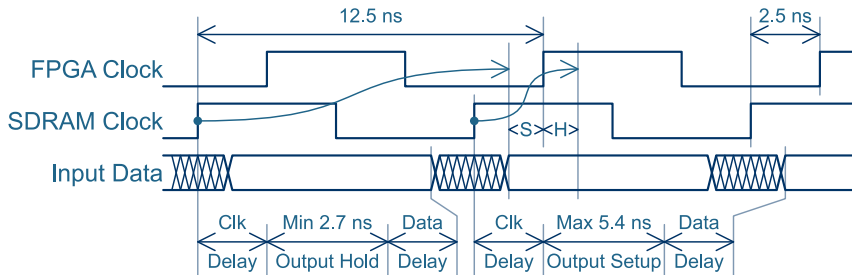


- ▶ Large FPGA internal delays and PCB trace delays
- ▶ Shift the external clock to ease the setup margin
- ▶ Multi-cycle requirement on the setup path (otherwise Quartus uses the 2.5 ns path) – the minimum propagation delay of 2.7 ns makes the 2.5 ns clock shift safe



# External Timing – Input


4 of 40



- ▶ Large FPGA internal delays and PCB trace delays
- ▶ Shift the external clock to ease the setup margin
- ▶ Multi-cycle requirement on the setup path (otherwise Quartus uses the 2.5 ns path) – the minimum propagation delay of 2.7 ns makes the 2.5 ns clock shift safe

# External Timing – Input


4 of 40



- ▶ Large FPGA internal delays and PCB trace delays
- ▶ Shift the external clock to ease the setup margin
- ▶ Multi-cycle requirement on the setup path (otherwise Quartus uses the 2.5 ns path) – the minimum propagation delay of 2.7 ns makes the 2.5 ns clock shift safe

# External Timing – Input

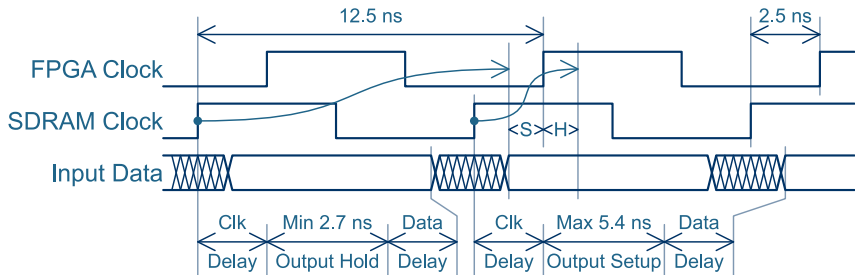
4 of 40



```
set_multicycle_path \
  -from [get_clocks {DRAM_CLK}] \
  -to   [get_clocks {*altpll_0*clk[0]}] \
  -setup 2
```

# External Timing – Input

4 of 40

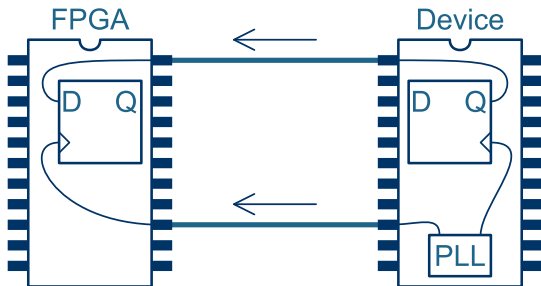


```
# Suppose 100 ps uncertainty and 200 ps PCB delay (each way)
set_input_delay -max -clock DRAM_CLK 5.9 [get_ports bpDRAM*]
set_input_delay -min -clock DRAM_CLK 3.0 [get_ports bpDRAM*]
```



# External Timing – DDR ADC

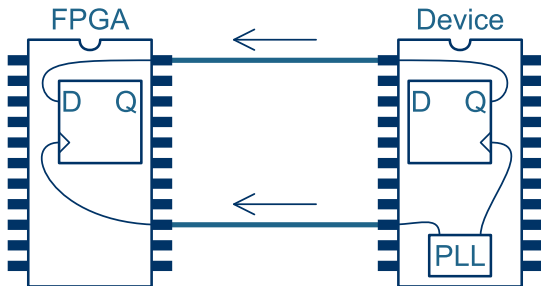
5 of 40



- ▶ The clock is sourced by the external device
- ▶ The PCB trace delays cancel
- ▶ The data is centre-aligned
- ▶ Quartus automatically does input port alignment;  
Xilinx must be manually tuned: use the IDELAY primitive

# External Timing – DDR ADC


5 of 40



- ▶ The clock is sourced by the external device
- ▶ The PCB trace delays cancel
- ▶ The data is centre-aligned
- ▶ Quartus automatically does input port alignment; Xilinx must be manually tuned: use the IDELAY primitive

# External Timing – DDR ADC

5 of 40




- ▶ The clock is sourced by the external device
- ▶ The PCB trace delays cancel
- ▶ The data is centre-aligned
- ▶ Quartus automatically does input port alignment;  
Xilinx must be manually tuned: use the IDELAY primitive



# External Timing – DDR ADC

5 of 40

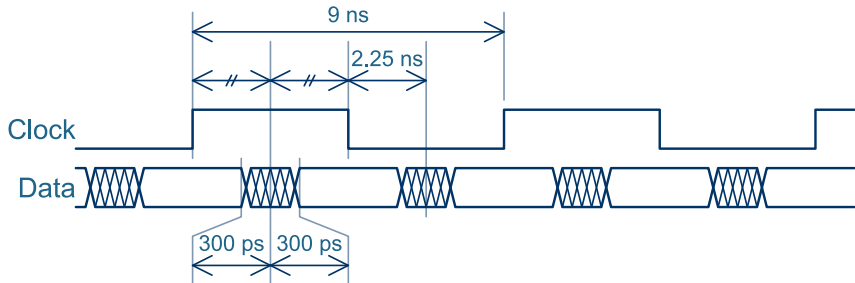


- ▶ The clock is sourced by the external device
- ▶ The PCB trace delays cancel
- ▶ The data is centre-aligned
- ▶ Quartus automatically does input port alignment; Xilinx must be manually tuned: use the IDELAY primitive



# External Timing – DDR ADC


5 of 40



```
# Suppose 100 ps uncertainty
create_clock      -name  ADC_DCO \
  -period 9 [get_ports ipADC_DCLK_P]
set_input_delay -clock ADC_DCO \
  -min 1.85 [get_ports ipADC_CH*]
set_input_delay -clock ADC_DCO \
  -max 2.65 [get_ports ipADC_CH*]
```

# External Timing – DDR ADC

5 of 40



```
set_input_delay -clock ADC_DCO -clock_fall \  
-min 1.85 [get_ports ipADC_CH*] -add_delay  
set_input_delay -clock ADC_DCO -clock_fall \  
-max 2.65 [get_ports ipADC_CH*] -add_delay
```

- ▶ Use the correct I/O standard
- ▶ Set the correct output current
- ▶ Set the pin capacitance

```
IOBUF PORT "opDRAM*" IO_TYPE=LVC MOS33 ;  
IOBUF PORT "bpDRAM*" IO_TYPE=LVC MOS33 ;  
IOBUF PORT "clkDRAM" IO_TYPE=LVC MOS33 ;
```



- ▶ Use the correct I/O standard
- ▶ Set the correct output current
- ▶ Set the pin capacitance

```
IOBUF PORT "opDRAM*" IO_TYPE=LVC MOS33 DRIVE=8 ;  
IOBUF PORT "bpDRAM*" IO_TYPE=LVC MOS33 DRIVE=8 ;  
IOBUF PORT "clkDRAM" IO_TYPE=LVC MOS33 DRIVE=8 ;
```



# Timing-related Pin Attributes

6 of 40

- ▶ Use the correct I/O standard
- ▶ Set the correct output current
- ▶ Set the pin capacitance

```
OUTPUT PORT "opDRAM*"      LOAD 3.8 pF ;  
OUTPUT PORT "bpDRAM_DQ*"   LOAD 6.0 pF ;  
OUTPUT PORT "CLK_DRAM"     LOAD 3.5 pF ;
```



# Outline

---

Advanced Timing Constraints

FIR Filters

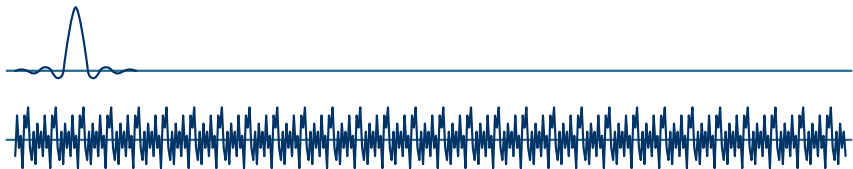
Practical – FIR Filter

Projects

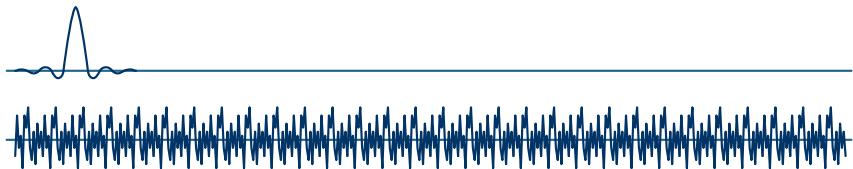
Tips and Tricks

Conclusion



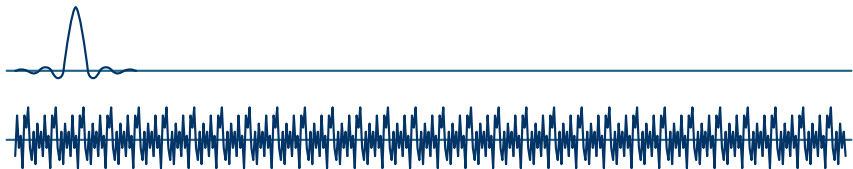


- Convolve the impulse-response with the signal:
  1. Time-reverse the impulse-response
  2. Within the FIR filter window, multiply the impulse-response sample by the signal sample
  3. Sum the products and output the result
  4. Move the impulse-response by one sample
  5. Repeat from step 2

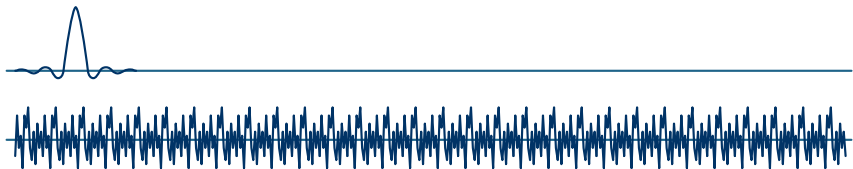


- Convolve the impulse-response with the signal:
  1. Time-reverse the impulse-response
  2. Within the FIR filter window, multiply the impulse-response sample by the signal sample
  3. Sum the products and output the result
  4. Move the impulse-response by one sample
  5. Repeat from step 2

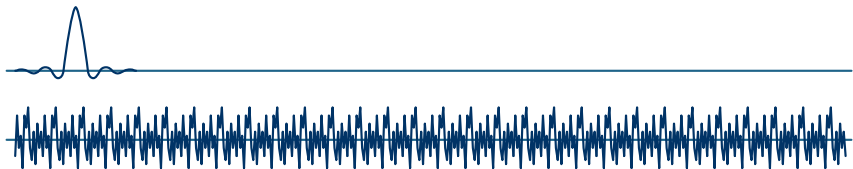




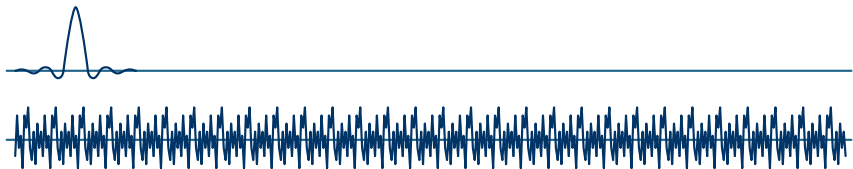
- Convolve the impulse-response with the signal:
  1. Time-reverse the impulse-response
  2. Within the FIR filter window, multiply the impulse-response sample by the signal sample
  3. Sum the products and output the result
  4. Move the impulse-response by one sample
  5. Repeat from step 2



- Convolve the impulse-response with the signal:
  1. Time-reverse the impulse-response
  2. Within the FIR filter window, multiply the impulse-response sample by the signal sample
  3. Sum the products and output the result
  4. Move the impulse-response by one sample
  5. Repeat from step 2



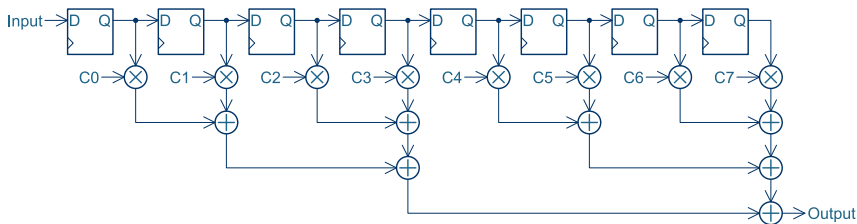
- ▶ Convolve the impulse-response with the signal:
  1. Time-reverse the impulse-response
  2. Within the FIR filter window, multiply the impulse-response sample by the signal sample
  3. Sum the products and output the result
  4. Move the impulse-response by one sample
  5. Repeat from step 2



- Convolve the impulse-response with the signal:
  1. Time-reverse the impulse-response
  2. Within the FIR filter window, multiply the impulse-response sample by the signal sample
  3. Sum the products and output the result
  4. Move the impulse-response by one sample
  5. Repeat from step 2

# FIR Filter – Implementation

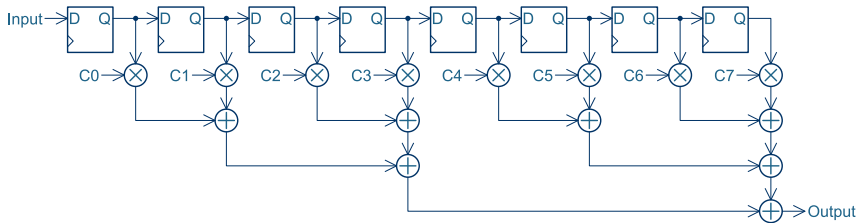
8 of 40



- Move the signal instead of the impulse-response
- To check the direction, inject an impulse (which should produce the impulse-response at the output)
- Pipeline the adder tree to increase the maximum clock frequency

# FIR Filter – Implementation

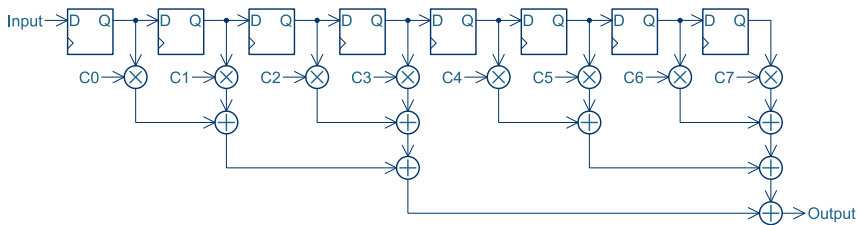
8 of 40



- ▶ Move the signal instead of the impulse-response
- ▶ To check the direction, inject an impulse (which should produce the impulse-response at the output)
- ▶ Pipeline the adder tree to increase the maximum clock frequency

# FIR Filter – Implementation

8 of 40



- ▶ Move the signal instead of the impulse-response
- ▶ To check the direction, inject an impulse (which should produce the impulse-response at the output)
- ▶ Pipeline the adder tree to increase the maximum clock frequency

- ▶ Always take the design criteria into account when designing an architecture
- ▶ What happens when you add decimation (i.e. you don't need an output sample every clock cycle)?
- ▶ What if the FPGA clock is much faster than the sample rate?
- ▶ What about a combination of the above?
- ▶ Always consider the scenario: sample rate, clock speed, power requirements, throughput requirements, decimation (if any), available resources, etc...





- ▶ Always take the design criteria into account when designing an architecture
- ▶ What happens when you add decimation (i.e. you don't need an output sample every clock cycle)?
- ▶ What if the FPGA clock is much faster than the sample rate?
- ▶ What about a combination of the above?
- ▶ Always consider the scenario: sample rate, clock speed, power requirements, throughput requirements, decimation (if any), available resources, etc...



- ▶ Always take the design criteria into account when designing an architecture
- ▶ What happens when you add decimation (i.e. you don't need an output sample every clock cycle)?
- ▶ What if the FPGA clock is much faster than the sample rate?
- ▶ What about a combination of the above?
- ▶ Always consider the scenario: sample rate, clock speed, power requirements, throughput requirements, decimation (if any), available resources, etc...

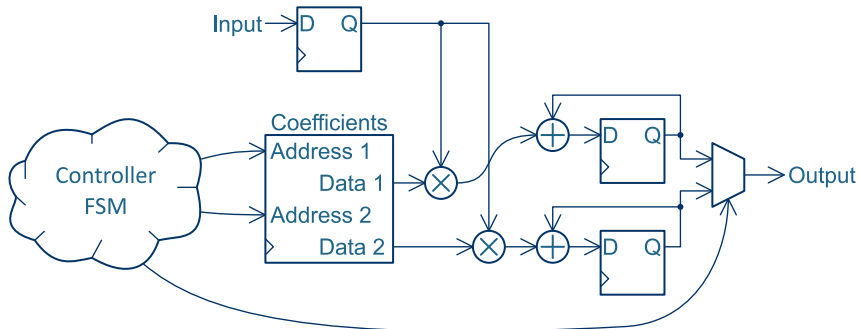


- ▶ Always take the design criteria into account when designing an architecture
- ▶ What happens when you add decimation (i.e. you don't need an output sample every clock cycle)?
- ▶ What if the FPGA clock is much faster than the sample rate?
- ▶ What about a combination of the above?
- ▶ Always consider the scenario: sample rate, clock speed, power requirements, throughput requirements, decimation (if any), available resources, etc...

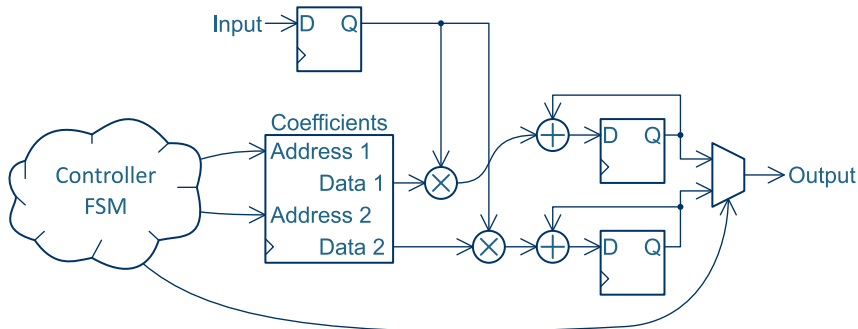


- ▶ Always take the design criteria into account when designing an architecture
- ▶ What happens when you add decimation (i.e. you don't need an output sample every clock cycle)?
- ▶ What if the FPGA clock is much faster than the sample rate?
- ▶ What about a combination of the above?
- ▶ Always consider the scenario: sample rate, clock speed, power requirements, throughput requirements, decimation (if any), available resources, etc...

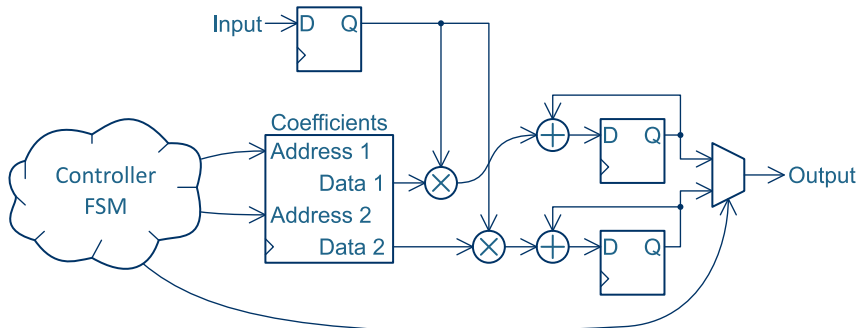




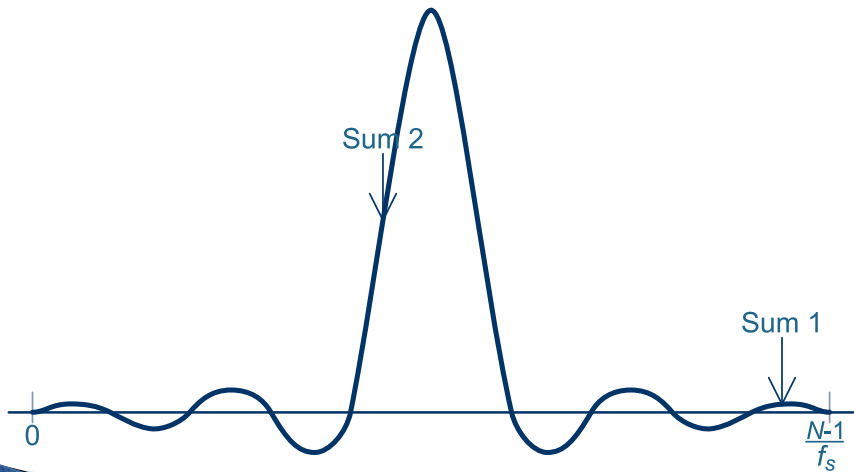
- ▶ This example decimates by  $N/2$ :  
a 512-point filter will decimate by 256
- ▶ The coefficient addresses are run  $N/2$  out of phase



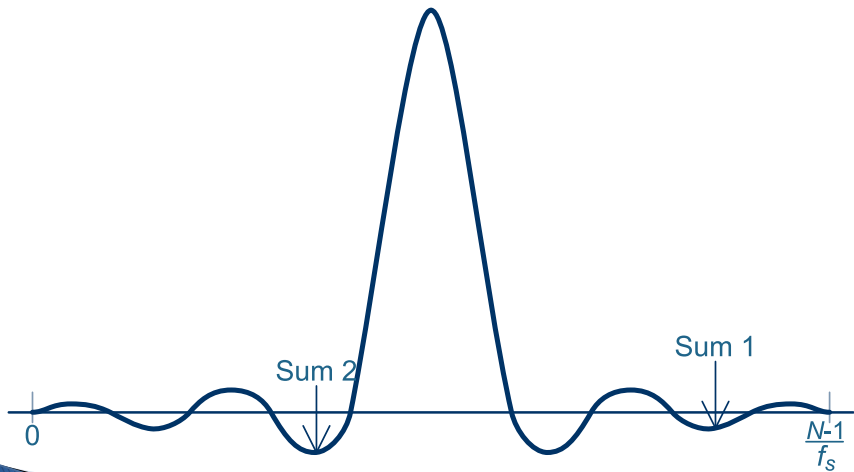
- This example decimates by  $N/2$ :  
a 512-point filter will decimate by 256
- The coefficient addresses are run  $N/2$  out of phase

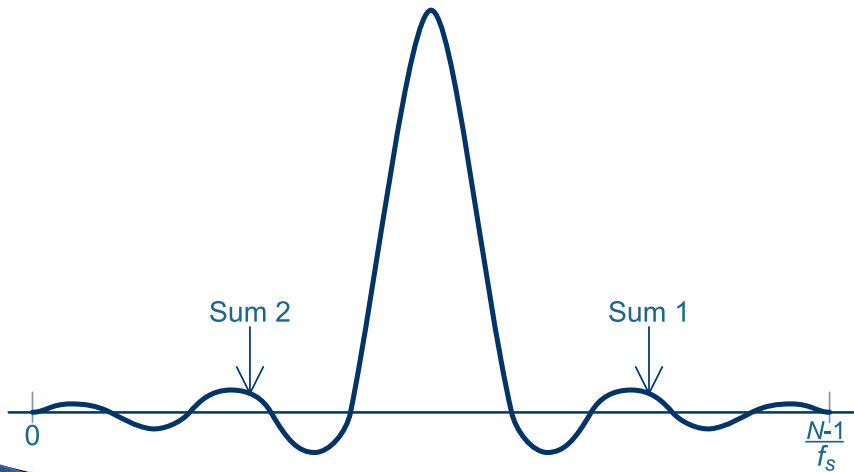


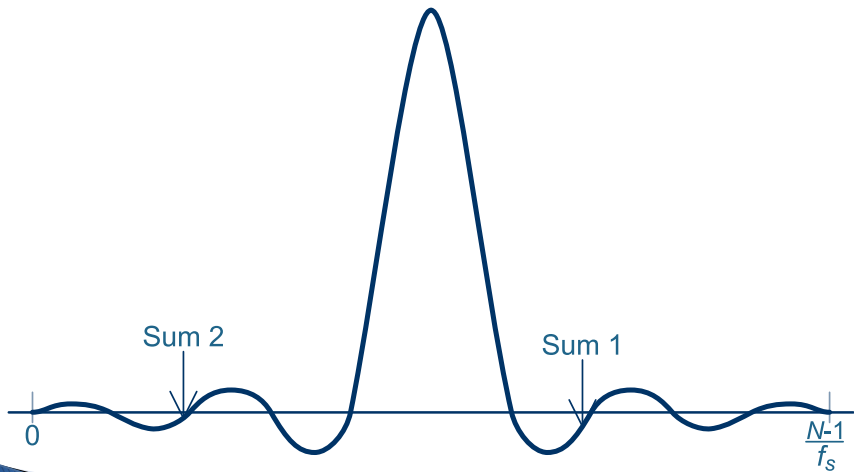
- ▶ This example decimates by  $N/2$ :  
a 512-point filter will decimate by 256
- ▶ The coefficient addresses are run  $N/2$  out of phase

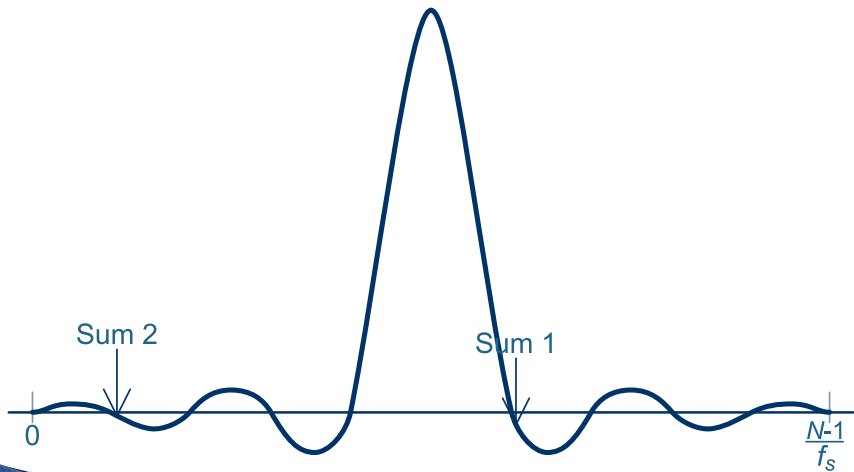


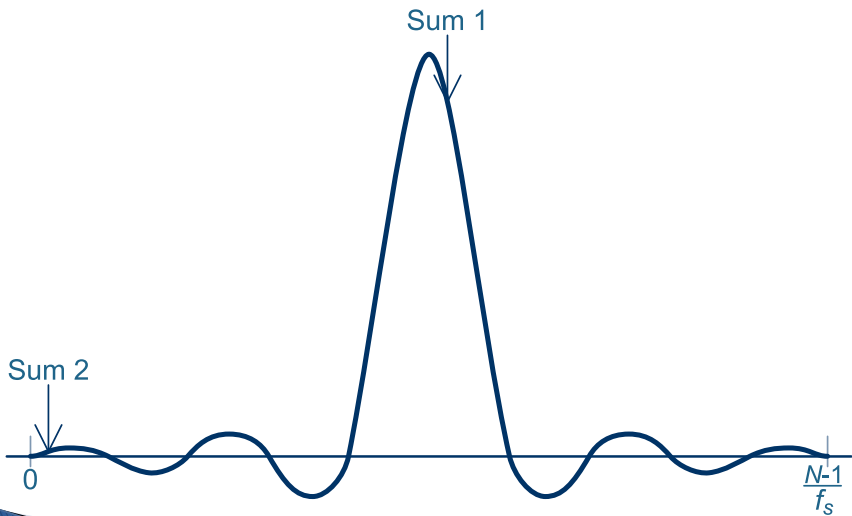


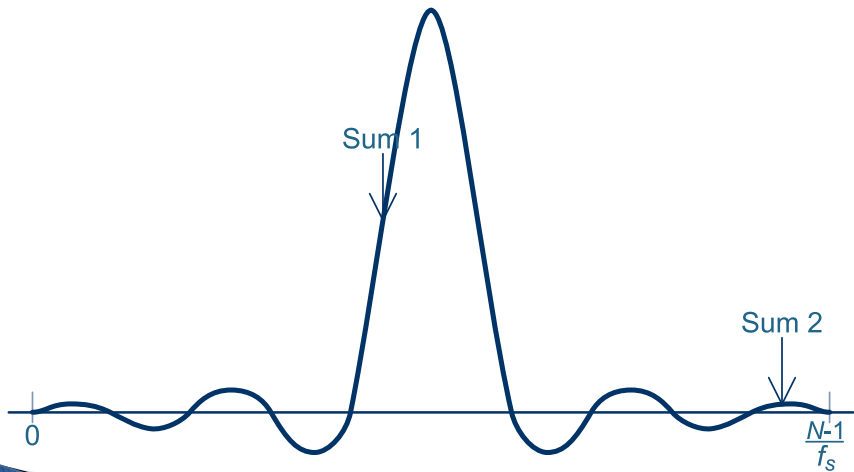


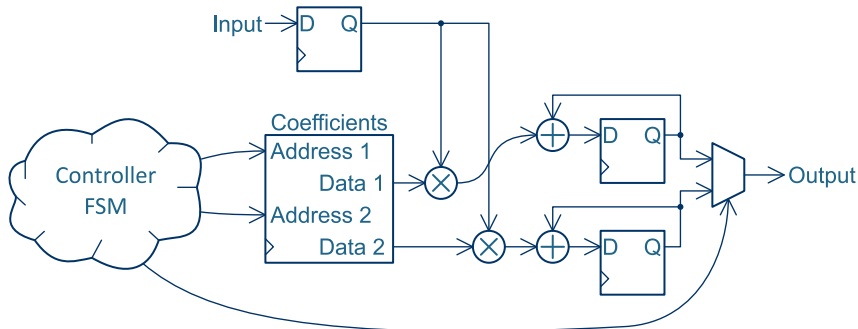













- If the sample clock is lower than the system clock, this same architecture can decimate by less:
- Keep more than 2 sums – sometimes more efficient to keep these in BRAM as well

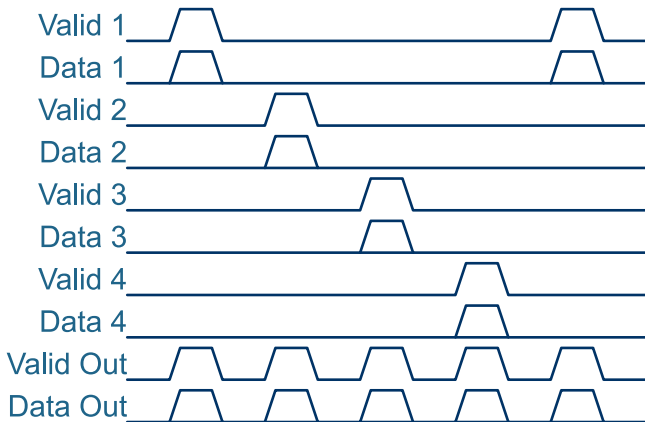



- ▶ If the sample clock is lower than the system clock, this same architecture can decimate by less:
- ▶ Keep more than 2 sums – sometimes more efficient to keep these in BRAM as well



# Combining FIR Filter Units

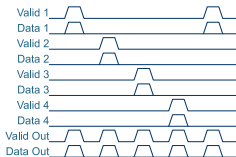
12 of 40



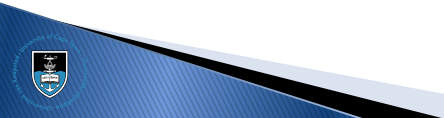


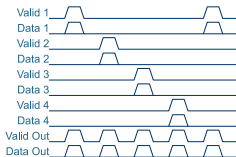
- ▶ Use multiple instances of a filter that decimates more than desired
- ▶ Reset the counters of the units out of phase
- ▶ Combine the outputs with a simple AND-OR circuit






- ▶ Use multiple instances of a filter that decimates more than desired
- ▶ Reset the counters of the units out of phase
- ▶ Combine the outputs with a simple AND-OR circuit



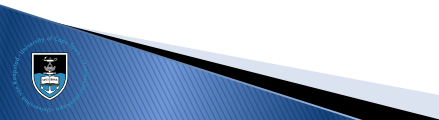


- ▶ Use multiple instances of a filter that decimates more than desired
- ▶ Reset the counters of the units out of phase
- ▶ Combine the outputs with a simple AND-OR circuit



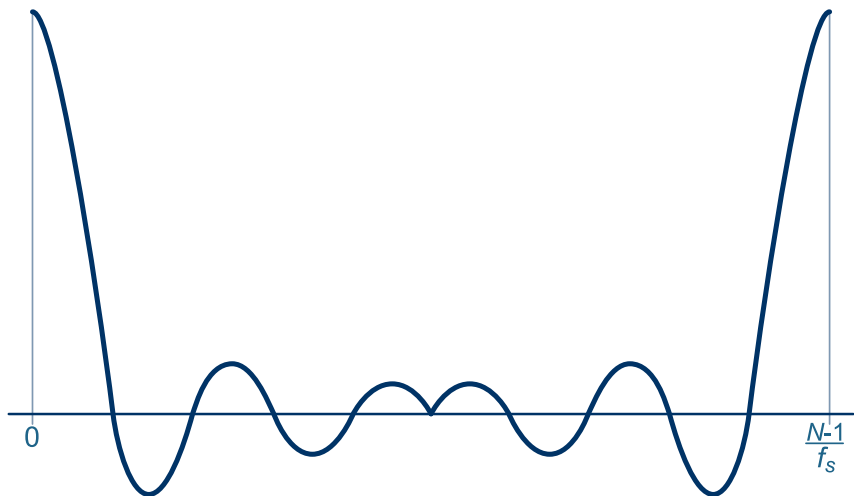
# FIR Filter – Desired Response

13 of 40



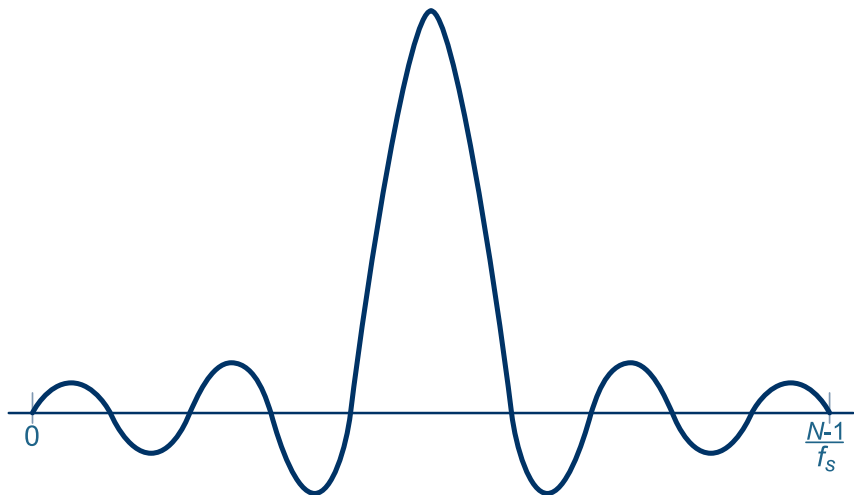
# FIR Filter – IFFT (real part)

13 of 40




# FIR Filter – Time-shifted

13 of 40



# FIR Filter – Windowed

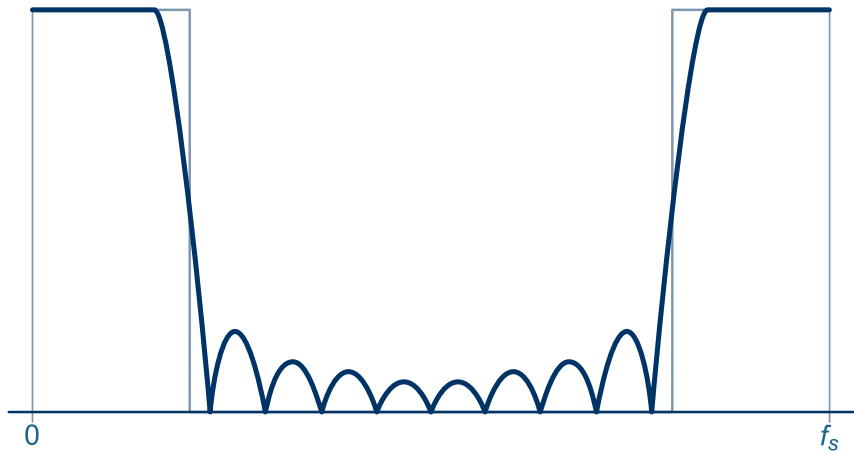
13 of 40






# FIR Filter – Actual Response

13 of 40



(Zero-pad to see the side-lobes)



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
- ▶ Verify on FPGA
- ▶ Combine eight filter units to drop the sub-sampling rate to N/8
- ▶ Verify on FPGA



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
- ▶ Verify on FPGA
- ▶ Combine eight filter units to drop the sub-sampling rate to N/8
- ▶ Verify on FPGA



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
- ▶ Verify on FPGA
- ▶ Combine eight filter units to drop the sub-sampling rate to N/8
- ▶ Verify on FPGA



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
  - ▶ Verify on FPGA
  - ▶ Combine eight filter units to drop the sub-sampling rate to N/8
  - ▶ Verify on FPGA



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
- ▶ Verify on FPGA
- ▶ Combine eight filter units to drop the sub-sampling rate to N/8
- ▶ Verify on FPGA



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
- ▶ Verify on FPGA
- ▶ Combine eight filter units to drop the sub-sampling rate to N/8
- ▶ Verify on FPGA



- ▶ Implement a N-point FIR filter that decimates by N (i.e. one sample output for every N samples input)
- ▶ Use an appropriate cut-off frequency and a **Hann window**  
$$w(n) = \sin^2 \left( \frac{\pi n}{N-1} \right)$$
- ▶ Use Matlab / Octave / Python to generate the FIR filter constants and memory initialisation file
- ▶ Verify through simulation
- ▶ Verify on FPGA
- ▶ Combine eight filter units to drop the sub-sampling rate to N/8
- ▶ Verify on FPGA





# Outline

---

Advanced Timing Constraints

FIR Filters

Practical – FIR Filter

Projects

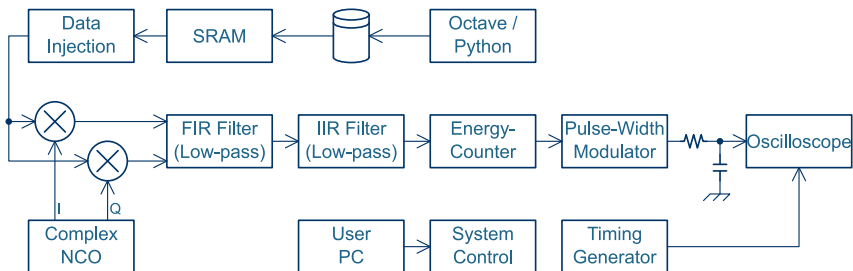
Tips and Tricks


Conclusion



# Practical – FIR Filter

15 of 40





# Coffee Break...

---

16 of 40



# Outline

---

Advanced Timing Constraints

FIR Filters

Practical – FIR Filter

Projects

Tips and Tricks

Conclusion



- ▶ Everybody must do a different project, but the projects are interlinked. Too make it more fun, make sure the system parameters are compatible across projects.
- ▶ You can propose a project: preferably in line with your current MSc research
- ▶ You need to design the DSP chain and choose appropriate system parameters
- ▶ Typically, a design will inject data from SDRAM into the DSP-chain, and then store the result in the same SDRAM, which is then read and analysed by the PC
- ▶ Your system must be controllable from the PC (typically over the UART)



- ▶ Everybody must do a different project
- ▶ You can propose a project: preferably in line with your current MSc research
- ▶ You need to design the DSP chain and choose appropriate system parameters
- ▶ Typically, a design will inject data from SDRAM into the DSP-chain, and then store the result in the same SDRAM, which is then read and analysed by the PC
- ▶ Your system must be controllable from the PC (typically over the UART)



- ▶ Everybody must do a different project
- ▶ You can propose a project: preferably in line with your current MSc research
- ▶ You need to design the DSP chain and choose appropriate system parameters: ADC sampling-rate, decimation (if any), architecture etc.
- ▶ Typically, a design will inject data from SDRAM into the DSP-chain, and then store the result in the same SDRAM, which is then read and analysed by the PC
- ▶ Your system must be controllable from the PC (typically over the UART)





- ▶ Everybody must do a different project
- ▶ You can propose a project: preferably in line with your current MSc research
- ▶ You need to design the DSP chain and choose appropriate system parameters: ADC sampling-rate, decimation (if any), architecture etc.
- ▶ Typically, a design will inject data from SDRAM into the DSP-chain, and then store the result in the same SDRAM, which is then read and analysed by the PC
- ▶ Your system must be controllable from the PC (typically over the UART)



- ▶ Everybody must do a different project
- ▶ You can propose a project: preferably in line with your current MSc research
- ▶ You need to design the DSP chain and choose appropriate system parameters: ADC sampling-rate, decimation (if any), architecture etc.
- ▶ Typically, a design will inject data from SDRAM into the DSP-chain, and then store the result in the same SDRAM, which is then read and analysed by the PC
- ▶ Your system must be controllable from the PC (typically over the UART)



- ▶ The demonstrations are scheduled for the week of 26 July – a date will be chosen at a later time
- ▶ You need to:
  - ▶ Demonstrate a working FPGA-based DSP chain
  - ▶ Give a 15-minute presentation on the design and performance results
  - ▶ Submit the source-code for review
  - ▶ Submit a project design report: nothing fancy – keep it to within 10 pages of overview, detail design and performance results.



- ▶ The demonstrations are scheduled for the week of 26 July – a date will be chosen at a later time
- ▶ You need to:
  - ▶ Demonstrate a working FPGA-based DSP chain
  - ▶ Give a 15-minute presentation on the design and performance results
  - ▶ Submit the source-code for review
  - ▶ Submit a project design report: nothing fancy – keep it to within 10 pages of overview, detail design and performance results.



- ▶ The demonstrations are scheduled for the week of 26 July – a date will be chosen at a later time
- ▶ You need to:
  - ▶ Demonstrate a working FPGA-based DSP chain
  - ▶ Give a 15-minute presentation on the design and performance results
  - ▶ Submit the source-code for review
  - ▶ Submit a project design report: nothing fancy – keep it to within 10 pages of overview, detail design and performance results.



- ▶ The demonstrations are scheduled for the week of 26 July – a date will be chosen at a later time
- ▶ You need to:
  - ▶ Demonstrate a working FPGA-based DSP chain
  - ▶ Give a 15-minute presentation on the design and performance results
  - ▶ Submit the source-code for review
  - ▶ Submit a project design report: nothing fancy – keep it to within 10 pages of overview, detail design and performance results.



- ▶ The demonstrations are scheduled for the week of 26 July – a date will be chosen at a later time
- ▶ You need to:
  - ▶ Demonstrate a working FPGA-based DSP chain
  - ▶ Give a 15-minute presentation on the design and performance results
  - ▶ Submit the source-code for review
  - ▶ Submit a project design report: nothing fancy – keep it to within 10 pages of overview, detail design and performance results.



- ▶ The demonstrations are scheduled for the week of 26 July – a date will be chosen at a later time
- ▶ You need to:
  - ▶ Demonstrate a working FPGA-based DSP chain
  - ▶ Give a 15-minute presentation on the design and performance results
  - ▶ Submit the source-code for review
  - ▶ Submit a project design report: nothing fancy – keep it to within 10 pages of overview, detail design and performance results.





- ▶ Remember that you are not designing the Radar / Communication system / etc.
- ▶ You are designing the FPGA-based processor, on a very small FPGA with limited resources
- ▶ Keep things simple – choose system parameters that favour easy implementation, not good system performance (for example: always assume that targets are slow-moving, that there are no multipath effects and that you have a low sampling-rate)
- ▶ At the same time, the project must show FPGA competence, so don't make it too trivial
- ▶ Use the tools available, including the libraries and high-level design tools, where appropriate



- ▶ Remember that you are not designing the Radar / Communication system / etc.
- ▶ You are designing the FPGA-based processor, on a very small FPGA with limited resources
- ▶ Keep things simple – choose system parameters that favour easy implementation, not good system performance (for example: always assume that targets are slow-moving, that there are no multipath effects and that you have a low sampling-rate)
- ▶ At the same time, the project must show FPGA competence, so don't make it too trivial
- ▶ Use the tools available, including the libraries and high-level design tools, where appropriate



- ▶ Remember that you are not designing the Radar / Communication system / etc.
- ▶ You are designing the FPGA-based processor, on a very small FPGA with limited resources
- ▶ Keep things simple – choose system parameters that favour easy implementation, not good system performance (for example: always assume that targets are slow-moving, that there are no multipath effects and that you have a low sampling-rate)
- ▶ At the same time, the project must show FPGA competence, so don't make it too trivial
- ▶ Use the tools available, including the libraries and high-level design tools, where appropriate



- ▶ Remember that you are not designing the Radar / Communication system / etc.
- ▶ You are designing the FPGA-based processor, on a very small FPGA with limited resources
- ▶ Keep things simple – choose system parameters that favour easy implementation, not good system performance (for example: always assume that targets are slow-moving, that there are no multipath effects and that you have a low sampling-rate)
- ▶ At the same time, the project must show FPGA competence, so don't make it too trivial
- ▶ Use the tools available, including the libraries and high-level design tools, where appropriate



- ▶ Remember that you are not designing the Radar / Communication system / etc.
- ▶ You are designing the FPGA-based processor, on a very small FPGA with limited resources
- ▶ Keep things simple – choose system parameters that favour easy implementation, not good system performance (for example: always assume that targets are slow-moving, that there are no multipath effects and that you have a low sampling-rate)
- ▶ At the same time, the project must show FPGA competence, so don't make it too trivial
- ▶ Use the tools available, including the libraries and high-level design tools, where appropriate



- ▶ I'm not on campus, but you can reach me via email
- ▶ Clearly explain what you're trying to do, and what you're struggling with
- ▶ Include your source code and, when applicable, some pictures of the architecture you're trying to implement and a test-bench to highlight the problem
- ▶ You'll find some good resources on Google, but be very careful – more often than not the people don't know what they're talking about and lead you down the wrong path



- ▶ I'm not on campus, but you can reach me via **email**
- ▶ Clearly explain what you're trying to do, and what you're struggling with
- ▶ Include your source code and, when applicable, some pictures of the architecture you're trying to implement and a test-bench to highlight the problem
- ▶ You'll find some good resources on Google, but be very careful – more often than not the people don't know what they're talking about and lead you down the wrong path



- ▶ I'm not on campus, but you can reach me via **email**
- ▶ Clearly explain what you're trying to do, and what you're struggling with
- ▶ Include your source code and, when applicable, some pictures of the architecture you're trying to implement and a test-bench to highlight the problem
- ▶ You'll find some good resources on Google, but be very careful – more often than not the people don't know what they're talking about and lead you down the wrong path







- ▶ I'm not on campus, but you can reach me via **email**
- ▶ Clearly explain what you're trying to do, and what you're struggling with
- ▶ Include your source code and, when applicable, some pictures of the architecture you're trying to implement and a test-bench to highlight the problem
- ▶ You'll find some good resources on Google, but be very careful – more often than not the people don't know what they're talking about and lead you down the wrong path



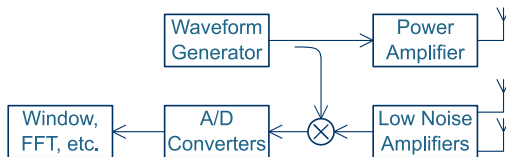





- ▶ **Sparse-array FMCW RADAR** (see Project 4)
- ▶ Choose system parameters appropriate for a practical radar – typical parameters include:
  - ▶ Sweep time of about 1 ms (sweep faster for fast-moving targets, and sweep slower for more range)
  - ▶ RF bandwidth of 500 MHz
  - ▶ 256 samples per sweep
  - ▶ 256 sweeps per burst (this is used for Doppler processing later in the chain)



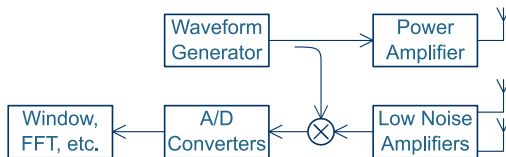
- ▶ Sparse-array FMCW RADAR (see Project 4)
- ▶ Choose system parameters appropriate for a practical radar – typical parameters include:
  - ▶ Sweep time of about 1 ms (sweep faster for fast-moving targets, and sweep slower for more range)
  - ▶ RF bandwidth of 500 MHz
  - ▶ 256 samples per sweep
  - ▶ 256 sweeps per burst (this is used for Doppler processing later in the chain)




- ▶ Sparse-array FMCW RADAR (see Project 4)
- ▶ Choose system parameters appropriate for a practical radar – typical parameters include:
  - ▶ Sweep time of about 1 ms (sweep faster for fast-moving targets, and sweep slower for more range)
  - ▶ RF bandwidth of 500 MHz
  - ▶ 256 samples per sweep
  - ▶ 256 sweeps per burst (this is used for Doppler processing later in the chain)



- ▶ Sparse-array FMCW RADAR (see Project 4)
- ▶ Choose system parameters appropriate for a practical radar – typical parameters include:
  - ▶ Sweep time of about 1 ms (sweep faster for fast-moving targets, and sweep slower for more range)
  - ▶ RF bandwidth of 500 MHz
  - ▶ 256 samples per sweep
  - ▶ 256 sweeps per burst (this is used for Doppler processing later in the chain)




- ▶ Sparse-array FMCW RADAR (see Project 4)
- ▶ Choose system parameters appropriate for a practical radar – typical parameters include:
  - ▶ Sweep time of about 1 ms (sweep faster for fast-moving targets, and sweep slower for more range)
  - ▶ RF bandwidth of 500 MHz
  - ▶ 256 samples per sweep
  - ▶ 256 sweeps per burst (this is used for Doppler processing later in the chain)




- ▶ Sparse-array FMCW RADAR (see Project 4)
- ▶ Choose system parameters appropriate for a practical radar – typical parameters include:
  - ▶ Sweep time of about 1 ms (sweep faster for fast-moving targets, and sweep slower for more range)
  - ▶ RF bandwidth of 500 MHz
  - ▶ 256 samples per sweep
  - ▶ 256 sweeps per burst (this is used for Doppler processing later in the chain)






- ▶ You can assume that the FPGA generates the transmit sweep triangle, so you can use the SDRAM address to determine where in the sweep you are
- ▶ Take the FFT of each sweep (range FFT), organise them into bursts and store the results in SDRAM
- ▶ This is the end of this project – another project could potentially take this output data and processes it further




- ▶ You can assume that the FPGA generates the transmit sweep triangle, so you can use the SDRAM address to determine where in the sweep you are
- ▶ Take the FFT of each sweep (range FFT), organise them into bursts and store the results in SDRAM
- ▶ This is the end of this project – another project could potentially take this output data and processes it further



- ▶ You can assume that the FPGA generates the transmit sweep triangle, so you can use the SDRAM address to determine where in the sweep you are
- ▶ Take the FFT of each sweep (range FFT), organise them into bursts and store the results in SDRAM
- ▶ This is the end of this project – another project could potentially take this output data and processes it further

# Project 2 – Doppler FFTs

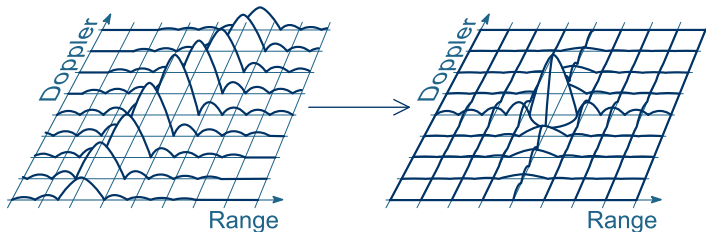
24 of 40



- ▶ Simulate the data output of Project 1 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and take the Doppler FFTs
- ▶ Store the resulting range-Doppler maps in SDRAM for the next step

# Project 2 – Doppler FFTs

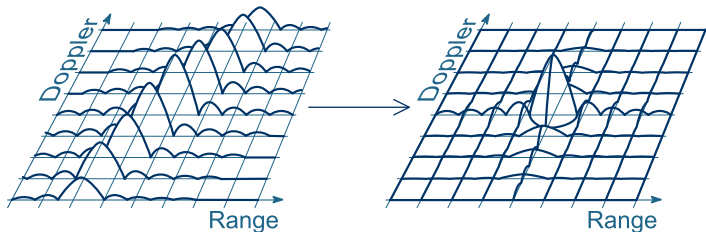
24 of 40



- ▶ Simulate the data output of Project 1 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and take the Doppler FFTs
- ▶ Store the resulting range-Doppler maps in SDRAM for the next step

# Project 2 – Doppler FFTs


24 of 40



- ▶ Simulate the data output of Project 1 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and take the Doppler FFTs
- ▶ Store the resulting range-Doppler maps in SDRAM for the next step

# Project 3 – Range CFAR

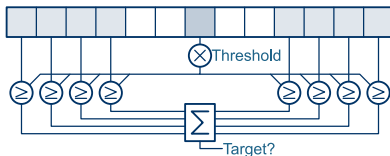
25 of 40



- ▶ Simulate the data output of Project 2 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and process the range CFAR
- ▶ Create a stream of detected targets (store the range, Doppler and phase of the two incoming channels)
- ▶ In a real system, this stream would go directly to the next step, but for this project, store the stream in SDRAM

# Project 3 – Range CFAR

25 of 40




- ▶ Simulate the data output of Project 2 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and process the range CFAR
- ▶ Create a stream of detected targets (store the range, Doppler and phase of the two incoming channels)
- ▶ In a real system, this stream would go directly to the next step, but for this project, store the stream in SDRAM



# Project 3 – Range CFAR

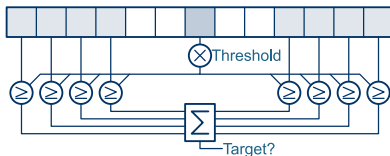
25 of 40



- ▶ Simulate the data output of Project 2 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and process the range CFAR
- ▶ Create a stream of detected targets (store the range, Doppler and phase of the two incoming channels)
- ▶ In a real system, this stream would go directly to the next step, but for this project, store the stream in SDRAM

# Project 3 – Range CFAR


25 of 40



- ▶ Simulate the data output of Project 2 and inject the anticipated result into the SDRAM
- ▶ Play back the corner-turned data and process the range CFAR
- ▶ Create a stream of detected targets (store the range, Doppler and phase of the two incoming channels)
- ▶ In a real system, this stream would go directly to the next step, but for this project, store the stream in SDRAM

# Project 4 – Angle Extraction


26 of 40



- ▶ Simulate the data output of Project 3 and inject the anticipated result into the SDRAM
- ▶ Play back the target stream and perform angle extraction for the sparse array

# Project 4 – Angle Extraction


26 of 40



- ▶ Simulate the data output of Project 3 and inject the anticipated result into the SDRAM
- ▶ Play back the target stream and perform angle extraction for the sparse array

# Project 5 – Pulsed Front-end

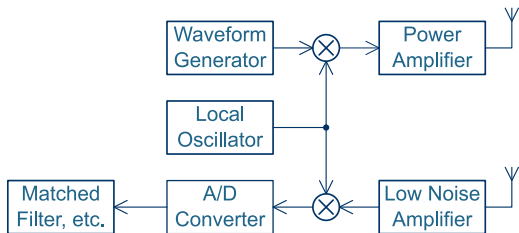
27 of 40



- Design and implement a chirped pulse RADAR front-end
- Inject raw ADC data and implement a matched filter by means of convolution (essentially a FIR filter)
- Display the PRI's on the oscilloscope

# Project 5 – Pulsed Front-end

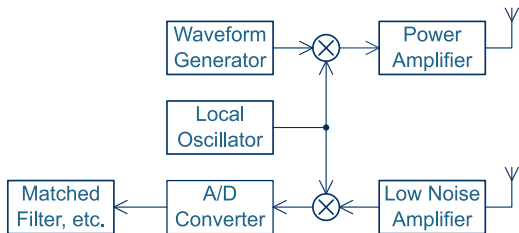
27 of 40



- ▶ Design and implement a chirped pulse RADAR front-end
- ▶ Inject raw ADC data and implement a matched filter by means of convolution (essentially a FIR filter)
- ▶ Display the PRI's on the oscilloscope

# Project 5 – Pulsed Front-end


27 of 40



- ▶ Design and implement a chirped pulse RADAR front-end
- ▶ Inject raw ADC data and implement a matched filter by means of convolution (essentially a FIR filter)
- ▶ Display the PRI's on the oscilloscope (Range is about  $6.7 \mu\text{s}/\text{km}$ , so it is practical to output the signal from a long-range RADAR (350 km or so) on 40 kHz bandwidth PWM (filter time-constant of  $4 \mu\text{s}$ )...

# Project 5 – Pulsed Front-end

27 of 40




- ▶ Design and implement a chirped pulse RADAR front-end
- ▶ Inject raw ADC data and implement a matched filter by means of convolution (essentially a FIR filter)
- ▶ Display the PRI's on the oscilloscope
- ▶ The rest of a typical processing chain is conceptually similar to projects 2 to 4, which can be adapted to process the results from this front-end



# Project 5 – Pulsed Front-end

28 of 40

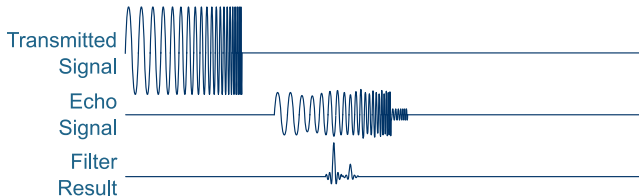


- ▶ You can assume that the FPGA generates the transmit timing control, so you can use the SDRAM address to determine where you are in the current PRI
- ▶ After doing matched filtering, organise the results into bursts and store them in SDRAM
- ▶ This is the end of this project – another project could potentially take this output data and processes it further



# Project 5 – Pulsed Front-end

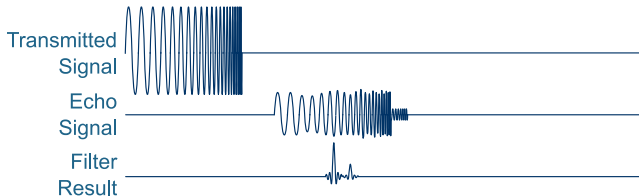
28 of 40



- ▶ You can assume that the FPGA generates the transmit timing control, so you can use the SDRAM address to determine where you are in the current PRI
- ▶ After doing matched filtering, organise the results into bursts and store them in SDRAM
- ▶ This is the end of this project – another project could potentially take this output data and processes it further

# Project 5 – Pulsed Front-end

28 of 40




- ▶ You can assume that the FPGA generates the transmit timing control, so you can use the SDRAM address to determine where you are in the current PRI
- ▶ After doing matched filtering, organise the results into bursts and store them in SDRAM
- ▶ This is the end of this project – another project could potentially take this output data and processes it further




# Project 6 – Commensal RADAR

29 of 40

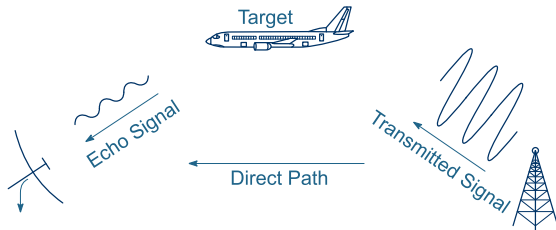


- ▶ Design and implement a commensal RADAR front-end
- ▶ Inject raw ADC data and implement range extraction (correlate the direct path with the echo signal)



# Project 6 – Commensal RADAR


29 of 40



- ▶ Design and implement a commensal RADAR front-end
- ▶ Inject raw ADC data and implement range extraction (correlate the direct path with the echo signal)

# Project 7 – FSK Communication

30 of 40

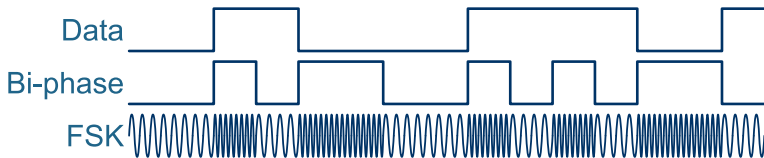


- ▶ Implement an FSK-based, bi-phase-coded communication channel
- ▶ Use S/PDIF as inspiration, with synchronisation word, etc.
- ▶ Inject simulated ADC data and demodulate by whatever means is convenient (matched filter, most likely)
- ▶ Display the received bit-stream on the oscilloscope
- ▶ Analyse channel performance in the presence of noise



# Project 7 – FSK Communication

30 of 40




- ▶ Implement an FSK-based, bi-phase-coded communication channel
- ▶ Use **S/PDIF** as inspiration, with synchronisation word, etc.
- ▶ Inject simulated ADC data and demodulate by whatever means is convenient (matched filter, most likely)
- ▶ Display the received bit-stream on the oscilloscope
- ▶ Analyse channel performance in the presence of noise



# Project 7 – FSK Communication

30 of 40



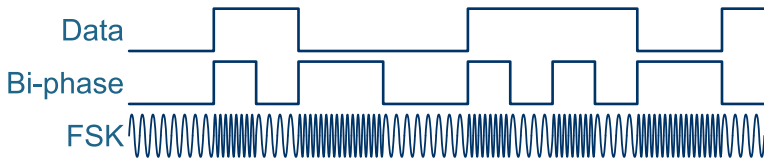
- ▶ Implement an FSK-based, bi-phase-coded communication channel
- ▶ Use **S/PDIF** as inspiration, with synchronisation word, etc.
- ▶ Inject simulated ADC data and demodulate by whatever means is convenient (matched filter, most likely)
- ▶ Display the received bit-stream on the oscilloscope
- ▶ Analyse channel performance in the presence of noise





# Project 7 – FSK Communication

30 of 40

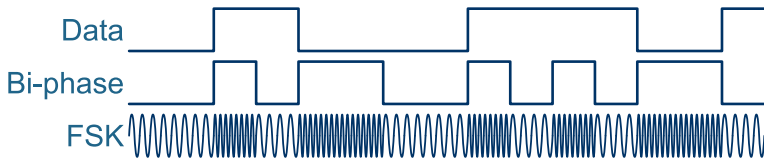


- ▶ Implement an FSK-based, bi-phase-coded communication channel
- ▶ Use **S/PDIF** as inspiration, with synchronisation word, etc.
- ▶ Inject simulated ADC data and demodulate by whatever means is convenient (matched filter, most likely)
- ▶ Display the received bit-stream on the oscilloscope
- ▶ Analyse channel performance in the presence of noise



# Project 7 – FSK Communication

30 of 40



- ▶ Implement an FSK-based, bi-phase-coded communication channel
- ▶ Use **S/PDIF** as inspiration, with synchronisation word, etc.
- ▶ Inject simulated ADC data and demodulate by whatever means is convenient (matched filter, most likely)
- ▶ Display the received bit-stream on the oscilloscope
- ▶ Analyse channel performance in the presence of noise





- ▶ Design and implement a 16-QAM modulator
- ▶ Inject a data stream and produce a 16-QAM output stream
- ▶ Including a synchronisation header and run-length limit
- ▶ Store the resulting modulated signal in SDRAM



- ▶ Design and implement a 16-QAM modulator
- ▶ Inject a data stream and produce a 16-QAM output stream
- ▶ Including a synchronisation header and run-length limit
- ▶ Store the resulting modulated signal in SDRAM



- ▶ Design and implement a 16-QAM modulator
- ▶ Inject a data stream and produce a 16-QAM output stream
- ▶ Including a synchronisation header and run-length limit
- ▶ Store the resulting modulated signal in SDRAM



- ▶ Design and implement a 16-QAM modulator
- ▶ Inject a data stream and produce a 16-QAM output stream
- ▶ Including a synchronisation header and run-length limit
- ▶ Store the resulting modulated signal in SDRAM



- ▶ Simulate the data output of Project 8 and inject the anticipated result into the SDRAM
- ▶ Assume an ideal RF front-end (i.e. no need to perform carrier-recovery)
- ▶ Play back the data stream and recover synchronisation
- ▶ Demodulate the data stream to obtain the original data

# Project 9 – QAM Demodulator

32 of 40



- ▶ Simulate the data output of Project 8 and inject the anticipated result into the SDRAM
- ▶ Assume an ideal RF front-end (i.e. no need to perform carrier-recovery)
- ▶ Play back the data stream and recover synchronisation
- ▶ Demodulate the data stream to obtain the original data



# Project 9 – QAM Demodulator

32 of 40



- ▶ Simulate the data output of Project 8 and inject the anticipated result into the SDRAM
- ▶ Assume an ideal RF front-end (i.e. no need to perform carrier-recovery)
- ▶ Play back the data stream and recover synchronisation
- ▶ Demodulate the data stream to obtain the original data

# Project 9 – QAM Demodulator


32 of 40



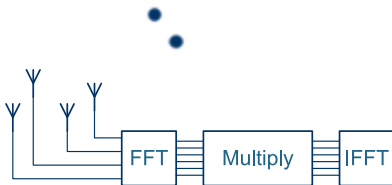
- ▶ Simulate the data output of Project 8 and inject the anticipated result into the SDRAM
- ▶ Assume an ideal RF front-end (i.e. no need to perform carrier-recovery)
- ▶ Play back the data stream and recover synchronisation
- ▶ Demodulate the data stream to obtain the original data

# Project 10 – Astronomy Receiver

33 of 40



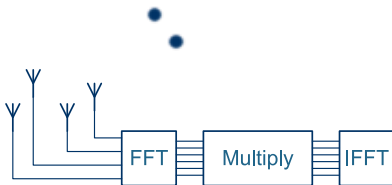
- ▶ Simulate a sky with at least two sources and inject ADC data from 4 receivers (4-bit per sample each)
- ▶ Perform correlation between all combinations of input channels (FFT, multiply, IFFT)
- ▶ Store the result in SDRAM
- ▶ Keep things simple: the earth is flat and stationary, etc.
- ▶ If parallel FFTs don't fit, do them sequentially



- ▶ Simulate a sky with at least two sources and inject ADC data from 4 receivers (4-bit per sample each)
- ▶ Perform correlation between all combinations of input channels (FFT, multiply, IFFT)
- ▶ Store the result in SDRAM
- ▶ Keep things simple: the earth is flat and stationary, etc.
- ▶ If parallel FFTs don't fit, do them sequentially

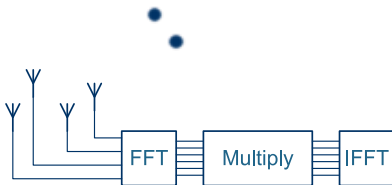
# Project 10 – Astronomy Receiver

33 of 40



- ▶ Simulate a sky with at least two sources and inject ADC data from 4 receivers (4-bit per sample each)
- ▶ Perform correlation between all combinations of input channels (FFT, multiply, IFFT)
- ▶ Store the result in SDRAM
- ▶ Keep things simple: the earth is flat and stationary, etc.
- ▶ If parallel FFTs don't fit, do them sequentially

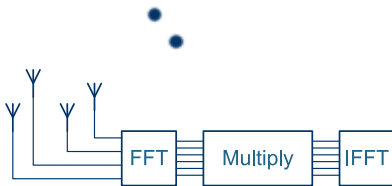




- ▶ Simulate a sky with at least two sources and inject ADC data from 4 receivers (4-bit per sample each)
- ▶ Perform correlation between all combinations of input channels (FFT, multiply, IFFT)
- ▶ Store the result in SDRAM
- ▶ Keep things simple: the earth is flat and stationary, etc.
- ▶ If parallel FFTs don't fit, do them sequentially

# Project 10 – Astronomy Receiver

33 of 40



- ▶ Simulate a sky with at least two sources and inject ADC data from 4 receivers (4-bit per sample each)
- ▶ Perform correlation between all combinations of input channels (FFT, multiply, IFFT)
- ▶ Store the result in SDRAM
- ▶ Keep things simple: the earth is flat and stationary, etc.
- ▶ If parallel FFTs don't fit, do them sequentially

# Project 11 – Astronomy Image

34 of 40



- ▶ Simulate the output of Project 10 and inject the anticipated result into the SDRAM
- ▶ Build an image from the correlation data
- ▶ Store the resulting image in SDRAM so that it can be viewed on the PC





# Project 11 – Astronomy Image

34 of 40



- ▶ Simulate the output of Project 10 and inject the anticipated result into the SDRAM
- ▶ Build an image from the correlation data
- ▶ Store the resulting image in SDRAM so that it can be viewed on the PC



# Project 11 – Astronomy Image

34 of 40



- ▶ Simulate the output of Project 10 and inject the anticipated result into the SDRAM
- ▶ Build an image from the correlation data
- ▶ Store the resulting image in SDRAM so that it can be viewed on the PC



# Outline

---

Advanced Timing Constraints

FIR Filters

Practical – FIR Filter

Projects

Tips and Tricks

Conclusion



- ▶ Whenever possible, design your modules such that they can be re-used in other projects
- ▶ Use module parametrisation when appropriate
- ▶ Use standardised bus structures and interfaces, and the same interface family across all projects
- ▶ Use consistent naming conventions
- ▶ Clearly mark negative logic in the name



- ▶ Whenever possible, design your modules such that they can be re-used in other projects
- ▶ Use module parametrisation when appropriate
- ▶ Use standardised bus structures and interfaces, and the same interface family across all projects
- ▶ Use consistent naming conventions
- ▶ Clearly mark negative logic in the name



- ▶ Whenever possible, design your modules such that they can be re-used in other projects
- ▶ Use module parametrisation when appropriate
- ▶ Use standardised bus structures and interfaces, and the same interface family across all projects
- ▶ Use consistent naming conventions
- ▶ Clearly mark negative logic in the name



- ▶ Whenever possible, design your modules such that they can be re-used in other projects
- ▶ Use module parametrisation when appropriate
- ▶ Use standardised bus structures and interfaces, and the same interface family across all projects
- ▶ Use consistent naming conventions
- ▶ Clearly mark negative logic in the name



- ▶ Whenever possible, design your modules such that they can be re-used in other projects
- ▶ Use module parametrisation when appropriate
- ▶ Use standardised bus structures and interfaces, and the same interface family across all projects
- ▶ Use consistent naming conventions
- ▶ Clearly mark negative logic in the name





- ▶ When possible, use the vendor-provided libraries
- ▶ Be careful: if the library function does not do quite what you want, it's generally easier and faster to design your own than to hack the existing library to do what you want
- ▶ Always use the correct tool for the job



- ▶ When possible, use the vendor-provided libraries
- ▶ Be careful: if the library function does not do quite what you want, it's generally easier and faster to design your own than to hack the existing library to do what you want
- ▶ Always use the correct tool for the job



- ▶ When possible, use the vendor-provided libraries
- ▶ Be careful: if the library function does not do quite what you want, it's generally easier and faster to design your own than to hack the existing library to do what you want
- ▶ Always use the correct tool for the job



- ▶ Always test as small a design as possible: in the ideal case, unit-test one module at a time
- ▶ In some cases, it's faster to simulate
- ▶ Other times, its easier and faster to compile and test on real hardware than to write the test-bench
- ▶ The on-chip logic analyser (i.e. Diamond Reveal, Quartus Signal-tap, Vivado Chip-scope, etc.) is your friend!



- ▶ Always test as small a design as possible: in the ideal case, unit-test one module at a time
- ▶ In some cases, it's faster to simulate
- ▶ Other times, its easier and faster to compile and test on real hardware than to write the test-bench
- ▶ The on-chip logic analyser (i.e. Diamond Reveal, Quartus Signal-tap, Vivado Chip-scope, etc.) is your friend!



- ▶ Always test as small a design as possible: in the ideal case, unit-test one module at a time
- ▶ In some cases, it's faster to simulate
- ▶ Other times, its easier and faster to compile and test on real hardware than to write the test-bench
- ▶ The on-chip logic analyser (i.e. Diamond Reveal, Quartus Signal-tap, Vivado Chip-scope, etc.) is your friend!



- ▶ Always test as small a design as possible: in the ideal case, unit-test one module at a time
- ▶ In some cases, it's faster to simulate
- ▶ Other times, its easier and faster to compile and test on real hardware than to write the test-bench
- ▶ The on-chip logic analyser (i.e. Diamond Reveal, Quartus Signal-tap, Vivado Chip-scope, etc.) is your friend!



- ▶ Most FPGA IDEs, including Lattice Diamond, Xilinx (AMD) Vivado and Altera (Intel) Quartus have powerful TCL scripting support
- ▶ You can automate version control, a more complicated compile chain (e.g. compiling a software component before the FPGA hardware), etc.
- ▶ Some IDEs lets you call a script automatically when starting and / or finishing a compile.





- ▶ Most FPGA IDEs, including Lattice Diamond, Xilinx (AMD) Vivado and Altera (Intel) Quartus have powerful TCL scripting support
- ▶ You can automate version control, a more complicated compile chain (e.g. compiling a software component before the FPGA hardware), etc.
- ▶ Some IDEs lets you call a script automatically when starting and / or finishing a compile.



- ▶ Most FPGA IDEs, including Lattice Diamond, Xilinx (AMD) Vivado and Altera (Intel) Quartus have powerful TCL scripting support
- ▶ You can automate version control, a more complicated compile chain (e.g. compiling a software component before the FPGA hardware), etc.
- ▶ Some IDEs lets you call a script automatically when starting and / or finishing a compile.



# Outline

---

Advanced Timing Constraints

FIR Filters

Practical – FIR Filter

Projects

Tips and Tricks

Conclusion



- ▶ We have covered a huge amount of work in a very short time
- ▶ It's up to you to go home and go play with the board
- ▶ And if you have questions: ask



- ▶ We have covered a huge amount of work in a very short time
- ▶ It's up to you to go home and go play with the board
- ▶ And if you have questions: ask



- ▶ We have covered a huge amount of work in a very short time
- ▶ It's up to you to go home and go play with the board
- ▶ And if you have questions: ask



# Select References

40 of 40



Stephen Brown and Zvonko Vranesic  
Fundamentals of Digital Logic with Verilog Design, 2<sup>nd</sup> Edition  
ISBN 978-0-07-721164-6



Merrill L Skolnik  
Introduction to RADAR Systems  
ISBN 978-0-07-288138-7



Mark A. Richards and James A. Scheer  
Principles of Modern Radar: Basic Principles  
ISBN 978-1-89-112152-4



Deepak Kumar Tala  
World of ASIC  
<http://www.asic-world.com/>



Jean P. Nicolle  
FPGA 4 Fun  
<http://www.fpga4fun.com/>



# FPGA Development for Radar, Radio-Astronomy and Communications

THE  
**RADAR**  
MASTERS COURSE



Dept. Electrical Engineering, University of Cape Town  
Private Bag, Rondebosch, 7701, South Africa  
<http://www.rrsg.uct.ac.za>



Presented by John-Philip Taylor  
Convened by Dr Stephen Paine

Day 6 – 4 May 2022